

Introduction to Cluster Computing Using Maxwell

Jerry Ebalunode

Center for Advanced Computation and Data Systems

(CACDS)

<http://cacds.uh.edu>



<http://support.cacds.uh.edu>

University of Houston

Houston, TX

@UHCACDS

facebook.com/CACDSATUH

Purpose of the Center

CACDS Purpose Statement

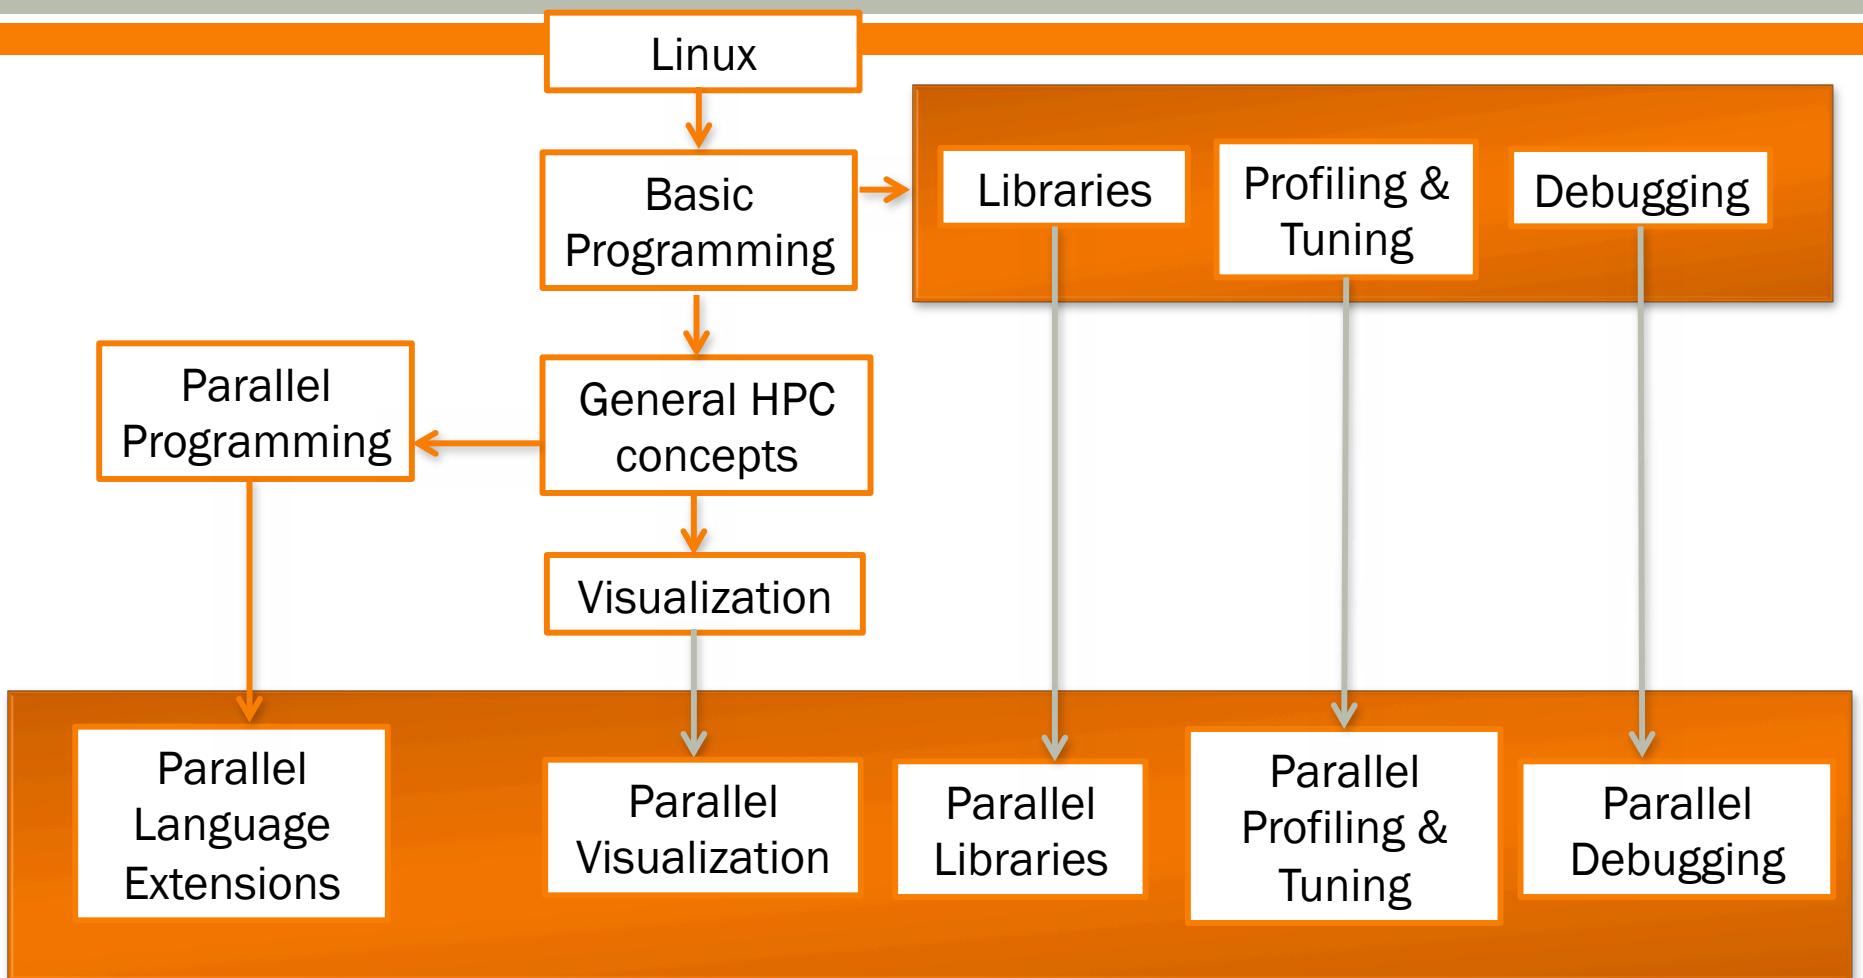
- Provide access to high performance computing and data storage resources
- Offer expert training, consulting, and technical support services
- Consult on research design and production
- Assist in the development and/or execution of supported research activities

Our Mission

- ∞ Advance the university's primary research thrusts in energy and health
- ∞ Support faculty research involving high performance computing (HPC), visualization, and computationally driven data-intensive science
- ∞ Provide consulting and training for UH faculty, staff and students
- ∞ Form relevant industry partnerships
- ∞ Facilitate collaborations between faculty as well as between faculty and external partners



HPC User Training Road Map



HPC Training

Fall 2014

Location: PGH 200

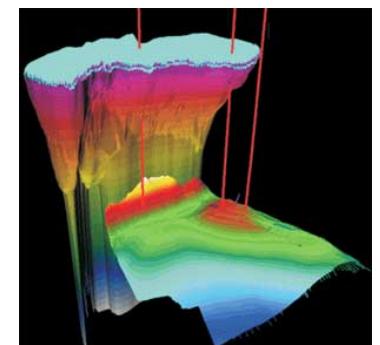
48 seat class room

- Linux
- Shell programming
- Cluster computing on Maxwell
- XSEDE Cyberinfrastructure and allocation process
- Python
- MATLAB
- C++
- Fortran
- Parallel Programming using OpenMP
- Parallel Programming using MPI
- Parallel numerical libraries

Location: PGH 235

12 seat class room

- Nvidia CUDA
 - introductory, intermediate and advanced
 - profiling, debugging & tuning
- OPENACC
 - intermediate, advanced
 - profiling, debugging & tuning
- Intermediate and Advanced MPI
 - profiling, debugging & tuning
- Intel Xeon Phi Programming
 - Native computing
 - OpenMP, Cilk++
 - Offload Execution
 - Symmetric Execution
 - profiling, debugging & tuning
- Visualization at scale
 - Paraview
 - VisIt



Training Registration

www.cacds.uh.edu/training

UNIVERSITYof **HOUSTON**

CENTER FOR ADVANCED COMPUTING & DATA SYSTEMS



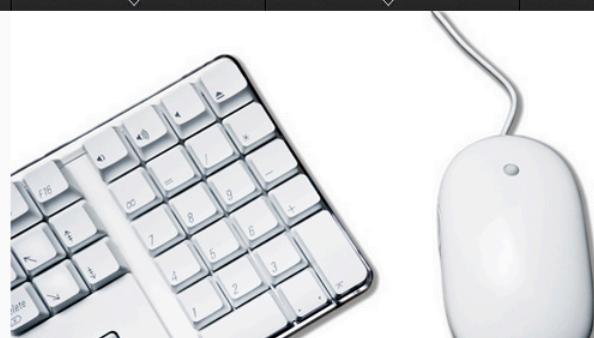
ABOUT

RESEARCH

EDUCATION

RESOURCES

NEWS



| Training Courses

EDUCATION

- Training Courses>>
- [Events](#)
- [Lectures](#)
- [Workshops](#)
- [Technical Seminars](#)

CACDS has launched several HPC training courses for members of the UH research community. Registration is required to take the courses. Please click the date you would like to attend under each course to register.

CURRENT TRAINING

CLASS NAME

- [Introduction to Linux](#)
- [Introduction to Shell Programming](#)
- [Introduction to Cluster Computing and Maxwell Cluster](#)
- [Introduction to C++ Programming](#)
- [Introduction to Python Programming](#)
- [Introduction to Parallel Programming with OpenMP](#)
- [GPGPU Parallel Programming with OpenACC](#)
- [GPGPU Parallel Programming with CUDA](#)
- [Introduction to Parallel Programming with MPI](#)
- [Introduction to Matlab](#)
- [Introduction to Intel Xeon Phi Programming](#)
- [Introduction to High Performance Numerical Libraries](#)
- [Visualization with VISIT I](#)
- [Visualization with VISIT II](#)
- [Visualization with Paraview I](#)

Overview

- ❑ Accessing the Cluster
- ❑ Maxwell Specifications
- ❑ User Environment
- ❑ Compiling Programs
- ❑ Portable Batch System (PBS)
- ❑ Batch Jobs
- ❑ Interactive Jobs
- ❑ Killing & Suspending Jobs

First Access Your Account

- ☞ Log into your accounts

- Username or login = hpc_user**X**
- Where **x** = sign in serial number 1 – 47
- **Password = cacds2014**
- Use your web browser
 - Firefox, Chromium or Google chrome

- ☞ Slides could be downloaded from URL below

/share/apps/tutorials/intro2maxwell.pdf

Disclaimer

- ∞ Maxwell cluster is for research work under a UH PI
 - Faculty sponsored RESEARCH activity only
- ∞ You will be using a surrogate cluster for this training in place of the Maxwell Cluster
 - You would not be using Maxwell for this training
- ∞ We assuming you are already familiar with the Linux Operating Environment
 - If not try to attend any of the upcoming Linux hands on tutorials offered by CACDS
- ∞ After this training, if you do not already have an account on Maxwell feel free to request one, after authorization from your PI(UH Faculty)
 - <http://www.rcc.uh.edu/cgi-bin/account.cgi> (actual application page)

Who Can Use Maxwell Cluster?

- ∞ Open to all UH faculty for research purposes
 - Students & postdocs – under UH faculty guidance
 - Collaborators – under UH faculty guidance

What Happens Next?

» How to connect to the Maxwell Cluster:

- For UNIX/Linux/Mac users
 - Open up terminal console and use Openssh secure shell client to connect
 - syntax:
 - **ssh username@cusco.hpcc.uh.edu**

Where **username** is your assigned login ID..

- you might be prompted to accept finger print...
 - Just type “yes” followed by Enter key
- For Windows users
 - Install Putty or any other compatible SSH client to connect
 - Putty URL = <http://www.putty.org/>

Connecting via OpenSSH Client on UNIX based Systems

```
bash-4.2$ ssh -Y jebaluno@cusco.hpcc.uh.edu  
The authenticity of host 'cusco.hpcc.uh.edu (129.7.54.9)' can't be established.
```

RSA key fingerprint is ac:3b:99:d0:6e:84:c0:d8:c9:16:60:d9:fe:e6:c7:71.

Are you sure you want to continue connecting (yes/no)? yes

Warning: Permanently added 'cusco.hpcc.uh.edu,129.7.54.9' (RSA) to the list of known hosts.

Use of University of Houston computing and network facilities requires prior authorization. Unauthorized access is prohibited. Usage may be subject to security testing and monitoring. Abuse is subject to criminal prosecution. A complete manual of security policies and procedures is available at <http://www.uh.edu/> in the Administration directory.

jebaluno@cusco.hpcc.uh.edu's password:

Connecting via OpenSSH Client on UNIX based Systems

jebaluno@cusco.hpcc.uh.edu's password:
Last login: Mon Oct 7 11:26:13 2013 from 129-7-249-234.dhcp.uh.edu

Maxwell Cluster

On My Computer

Sent Items

UH (University of Houston)

USC (University of South Carolina)

Deleted Items

On My Computer

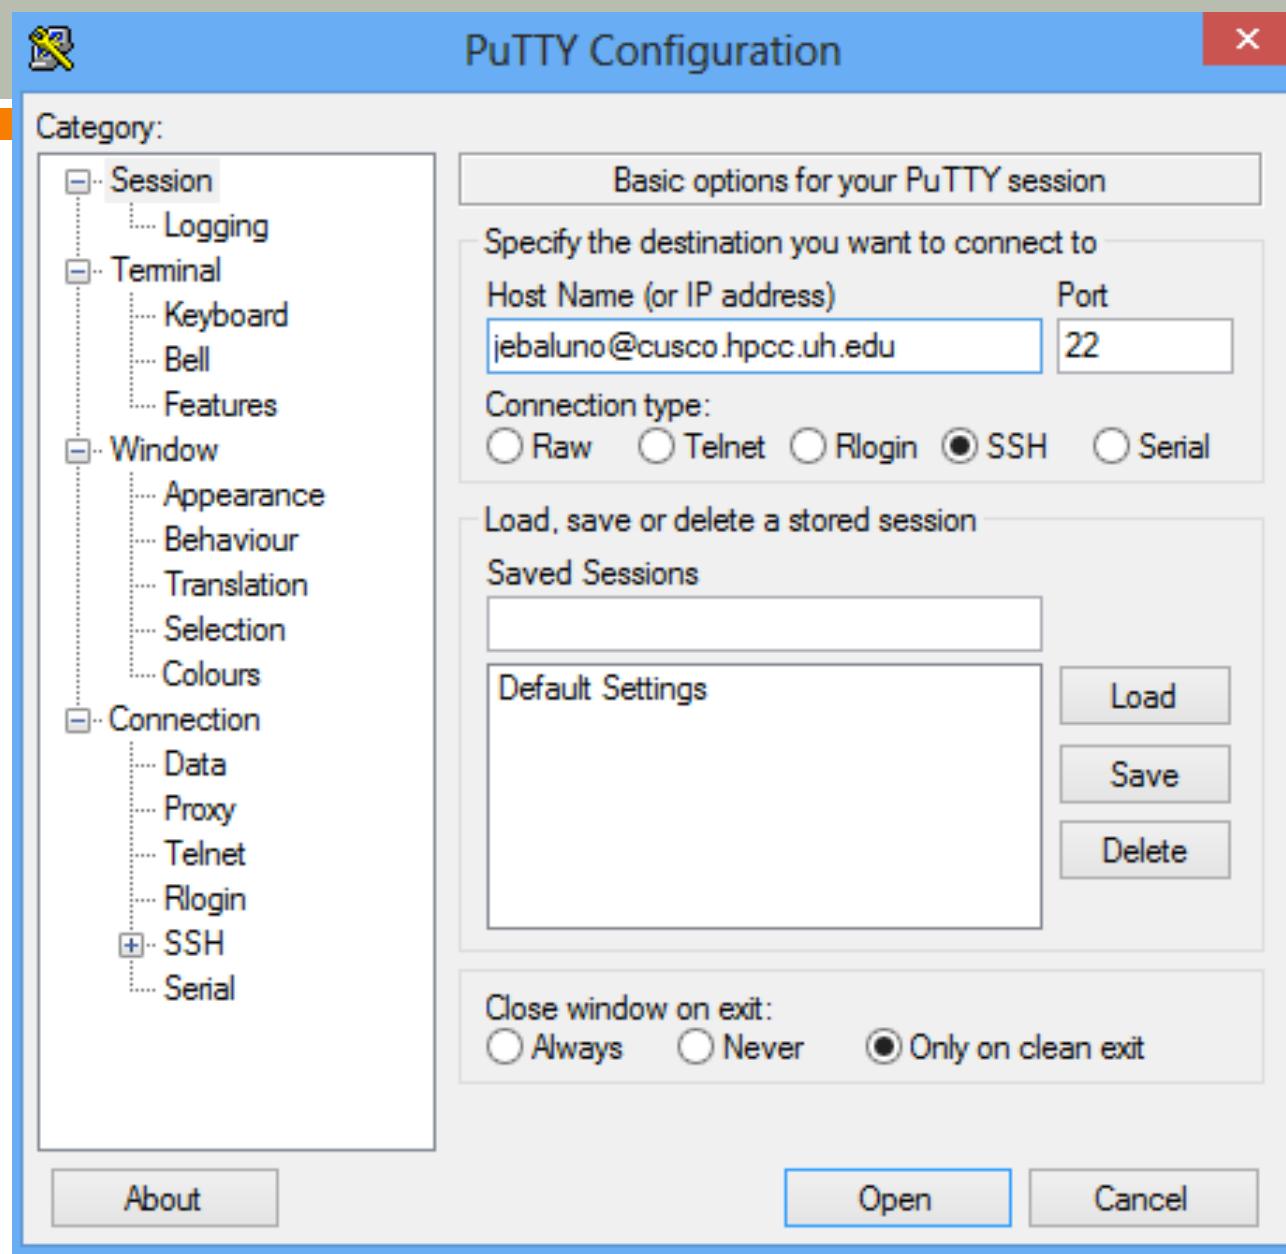
Conflicts

Operated by IT HPC Group | <http://www.rcc.uh.edu>
Help?? Email HPC | help@hpcc.uh.edu
Quick Reference Doc | </usr/local/doc/maxwellquickref.pdf>

[jebaluno@cusco ~]\$ hostname
cusco.hpcc.uh.edu
[jebaluno@cusco ~]\$

The screenshot shows a desktop interface with several windows open. In the foreground, a terminal window displays the command 'hostname' and its output 'cusco.hpcc.uh.edu'. Behind the terminal, there is a file browser window titled 'Maxwell Cluster' showing various email folders like 'Sent Items', 'UH (University of Houston)', and 'USC (University of South Carolina)'. A system tray or dock is visible at the bottom, containing icons for network status, battery level, and other system information.

Connecting via Putty on Windows Systems



Connecting via Putty on Windows Systems

```
Using username "jebaluno".
```

```
Use of University of Houston computing and network facilities requires  
prior authorization. Unauthorized access is prohibited. Usage may be  
subject to security testing and monitoring. Abuse is subject to  
criminal prosecution. A complete manual of security policies and  
procedures is available at http://www.uh.edu/ in the Administration  
directory.
```

```
jebaluno@cusco.hpcc.uh.edu's password: █
```

Note: you will not be able see your password

Accessing Tutorial Materials

First login into the cluster:

TYPE AND EXECUTE COMMANDS IN RED!!!

ssh -Y 10.1.1.1

cd

cp /share/apps/tutorials/intro2maxwell.zip ~

unzip intro2maxwell.zip

cd intro2maxwell

The Maxwell Cluster

- Both Shared and Distributed systems
 - single core, multicore (multithreaded) and MPI jobs
- 434+ compute nodes (3756+ cpu cores)
 - <http://baragon.hpcc.uh.edu/ganglia>
 - 22+ Geforce GTX 570 Nvidia GPUs
 - 8+ Nvidia Tesla C2075 GPUs
- Operating system: Linux
 - Distribution: Redhat 6.5
- Gigabit network interconnect
- Data Storage:
 - Default: 10+ GB Home directory space per user
 - PI's could request for project directory space to host files needed for simulations
 - Justification might be required

Specifications

Name	Count	Processor Type	Speed GHz	Total CPU Cores	Memory/ core	GPU
cusco (login node)	1	Quad-Core AMD Opteron(tm) Processor 2354	2.20	8	32+ GB	0.00
compute-X-Y	183	Quad-Core AMD Opteron(tm) Processor 2354	2.20	8	16+ GB	0.00
compute-X-Y	25	Quad-Core AMD Opteron(tm) Processor 8356	2.30	32	64+ GB	0.00
compute-X-Y	39	Quad-Core AMD Opteron(tm) Processor 2376	2.30	8	16.3+ GB	0.00
compute-X-Y	112	Dual Core AMD Opteron(tm) Processor 275	2.20	4	8+ GB	0.00
compute-X-Y	28	AMD Opteron(TM) Processor 6274	2.50	32	65+ GB	0.00
compute-X-Y	22	Intel(R) Xeon(R) CPU 5160	3.00	4	32+ GB	GeForce GTX 570
compute-X-Y	8	Intel(R) Xeon(R) CPU 5160	3.00	4	32+ GB	Tesla C2075

X = Rack Number

Y = Node Number

HPC Software Available on Maxwell

- ☞ Proprietary & Open-Source Software:
 - Examples:
 - Amber, Autodock, Bwa, NwChem, Espresso, Orca, R, Octave, Matlab, Gromacs, Lammmps, NAMD, Turbomole, VASP, Rosetta, R, NAG, MKL , ASE, GSL, OpenMPI, Intel, Gnu and PGI compilers
 - and much more...
 - Proprietary software may be limited to licensed users e.g. Fluent, VASP
 - Open-source software is accessible to all
- ☞ PI could request for installation of additional software
 - PI's licensed propriety application can be installed too
- ☞ Software environment management via “**modules**”
 - Allows software to be accessed via an assigned module
 - Dynamic modification of a user's environment

Computing Environment Setup Using Modules

- ☞ The module system is used to make software and related settings available easily
 - software environments can be loaded and unloaded dynamically
- ☞ To get a list of available software:
module avail
(also useful: **module list**)
- ☞ To clear-out all added modules:
module purge
- ☞ To clear-out all specific modules:
module rm module1 module2
- ☞ To use specific software every time you log-in, place the
module add command in your `~/.bashrc` file i.e:
echo "module add intel openmpi" >> ~/.bashrc
- ☞ Now add intel compiler and MPI runtime (Openmpi version) to your environment)
module add openmpi

Compiling Programs

- ❖ CPU programs (compiled to execute on central processing unit)
 - Serial Programs
 - Parallel Programs
- ❖ GPU programs (compiled to execute on graphics processing unit)
 - GPGPU Programs

Compiling Programs: CPU Serial Program

Example 1 Program in file: gethostname.c

```
#include <stdio.h>
#include <sys/utsname.h>
int main ( )
{
    struct utsname uts;
    uname (&uts);
    printf ("Process on node %s.\n",
    uts.nodename);
    return 0;
}
```

- To compile:
 - 1. choose compiler
module add intel
 - 2. Compile
icc gethostname.c -o gethostname.exe

Compiling Programs: CPU Parallel Program

Programs using Message Passing Interface (MPI)

☞ Example: 1 Program in file : gethostname.mpi.c

```
#include <stdio.h>
#include <sys/utsname.h>
#include <mpi.h>
int main (int argc, char *argv[]){
    struct utsname uts;
    int rank;
    MPI_Init (&argc, &argv);
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    uname (&uts);
    printf ("Process %d on node %s.\n", rank, uts.nodename);
    MPI_Finalize ();
    return 0;
}
```

- To compile:
 - 1. choose compiler
`module add openmpi`
 - 2. Compile
`mpicc gethostname.mpi.c -o gethostname.mpi.exe`

Compiling Programs: GPGPU Parallel Program

Example 1 Program in file: sortkeys_basic_thrust.cu

```
#include <thrust/host_vector.h>
#include <thrust/device_vector.h>
#include <thrust/sort.h>
int main(void)
{
    // generate 16 Million random numbers on the host
    thrust::host_vector<int> h_vec(1 << 24);
    thrust::generate(h_vec.begin(), h_vec.end(), rand);
    thrust::device_vector<int> d_vec= h_vec; // transfer data to the device
    thrust::sort(d_vec.begin(), d_vec.end()); // sort data on the device (805 Mkeys/sec on GeForce GTX 480)
    thrust::copy(d_vec.begin(), d_vec.end(), h_vec.begin()); // transfer data back to host
    return 0;
}
```

- To compile:
 - 1. choose compiler

```
module add cuda-toolkit
2. Compile name = nvcc
nvcc sortkeys_basic_thrust.cu -o sortkeys_basic_thrust.exe
```

The Queue System (PBS)



- ❖ Portable Batch System (aka PBS) is the resource management service used on the Maxwell Cluster
- ❖ Enables you to make more efficient use of your time through scripting computational tasks
- ❖ PBS takes care of running these tasks and returning the results
- ❖ If the cluster is full, PBS holds your tasks and runs them when the resources are available
- ❖ PBS ensures fair sharing of cluster resources (policy enforcement)
 - PBS ensures optimal/efficient use of available resources

PBS Commands

Command(s)	Description
qstat	Check status of all jobs
qsub ./myjob	Submit batch jobs
qstat -u \$USER	Check status of just your jobs
qsub -I	Submit an interactive job
qhold jobID i.e. qhold 12345	Put a job on hold (before it starts)
qrsl jobID	Release a job from hold status
qdel #jobID	Delete a job, running or not



Using qsub with a Job Script

- ❖ Syntax/Format:

- **qsub [qsub params] [job script file]**

- ❖ Examples:

- **qsub myjob**

PBS Batch Job Script File

- » Very simple example:
- » Serial job requesting 1 CPU, 2 GB memory, 60 hours of walltime

```
#!/bin/bash
#PBS -N jobname
#PBS -l mem=2gb,walltime=60:00:00
#PBS -j y
```

cd \$PBS_O_WORKDIR

./my.program.exe [arguments for my program]

Using Parallel Environment to Manage Memory and CPU Core Availability

- ☞ `#PBS -l nodes=16:ppn=1` --use only one CPU core per node and use a total of 16 processes.
- ☞ `#PBS -l nodes=8:ppn=2` --use only 2 CPU cores per node and use a total of 16 processes
- ☞ `#PBS -l nodes=4:ppn=4` --use only 4 CPU cores per node and use a total of 16 processes
- ☞ `#PBS -l nodes=8:ppn=8` --use only 8 CPU cores per node and use a total of 64 processes

qsub Notification Parameters

- ☞ **#PBS -M [e-mail address]**
 - e-mail address can be a list of email addresses
 - separated by commas.

- ☞ **#PBS -m bea or -m be or -m e**
 - send an email when the job **b**egins, **e**nds, or is **a**borted.

Example:

```
#!/bin/bash
#PBS -N jobname
#PBS -j y
#PBS -M jerry@uh.edu
#PBS -m abe
```

```
cd $PBS_O_WORKDIR
```

```
./my.program.exe [arguments for my program]
```

Sample Notification Email

PBS JOB 3687125.cusco.hpcc.uh.edu

root

Sent: Monday, October 7, 2013 at 6:27 AM

To: jebalunode@uh.edu

PBS Job Id: 3687125.cusco.hpcc.uh.edu

Job Name: waterint-ACC-H17

Exec host: compute-7-25/31+compute-7-25/30+compute-7-25/29+compute-7-25/28+compute-7-25/27+compute-7-25/26+compute-7-25/25+compute-7-25/24+compute-7-25/23+compute-7-25/22+compute-7-25/21+compute-7-25/20+compute-7-25/19+compute-7-25/18+compute-7-25/17+compute-7-25/16+compute-7-25/15+compute-7-25/14+compute-7-25/13+compute-7-25/12+compute-7-25/11+compute-7-25/10+compute-7-25/9+compute-7-25/8+compute-7-25/7+compute-7-25/6+compute-7-25/5+compute-7-25/4+compute-7-25/3+compute-7-25/2+compute-7-25/1+compute-7-25/0

Execution terminated

Exit_status=0

resources_used.cput=08:46:15

resources_used.mem=352472kb

resources_used.vmem=35667756kb

resources_used.walltime=00:33:07

Handling Output Files

- ❖ When a job is started it takes its job name from the script file that was submitted.
 - The standard output and error output are sent to files named jobname.o987349 and jobname.e987349 (in your job directory)
- ❖ The following parameters modify this behavior:
 - **-e** [path] standard error output file.
 - **-o** [path] standard output file
 - **-j y** merge the error and standard output
- ❖ **#PBS -N** jobname name to be used for the job.

Using qsub with a Job Script

- ❖ Syntax/Format:

- **qsub [qsub params] [script file]**

- ❖ More examples on using parameters at the command line:

- **qsub -N test1 myjob.sh**
 - **qsub -l nodes=2:ppn=2 -o test2.out -N test2 myjob.sh**
 - **qsub -l walltime=02:00:00,mem=4gb myjob.sh**
 - **qsub -l nodes=2:ppn=32,mem=64gb myjob.sh**
 - **qsub -q gpu -l nodes=1:ppn=4 my_gpu_job.sh**

- ❖ To see all available qsub options, just run **qsub --help**

Using an Interactive Job

- ∞ Special kind of batch job
- ∞ Useful for debugging applications, short tests, or for computational steering

The diagram illustrates the use of interactive job submission commands. It features two orange rounded rectangular boxes at the top: the left one is labeled "Capital ‘I’" and the right one is labeled "Small L ‘I’". Below these labels, several blue text lines represent command examples, each with an arrow pointing to its corresponding label. The commands are:

- qsub -I (points to Capital “I”)
- qsub -I -l (points to Small L “I”)
- qsub -I -l walltime=2:00:00 (points to Small L “I”)
- qsub -I -l walltime=2:00:00,nodes=2:ppn=8 (points to Small L “I”)
- qsub -I -l walltime=2:00:00 -q gpu (points to Small L “I”)

Job Queues

Queue Name	Description
short	Queue to run short jobs (=< 4 hours)
medium	Queue to run medium jobs (1 week allowed)
long	Queue to run very long jobs (2 weeks allowed)
gpu	Queue to run gpu jobs on GTX 570 equipped nodes (2 weeks allowed)
gpu-tesla	Queue to run gpu jobs on Testla C2075 equipped nodes (2 weeks allowed)

USE “qstat -q” to probe the queues installed

Note gpu-tesla and gpu queue are not available/functional on the surrogate cluster

PBS Script for Serial Program

1 CPU

🔗 Example of a Serial Job Script

filename: job.gethostname.bash

```
#!/bin/bash
#PBS -N gethost
#PBS -l mem=8gb,walltime=00:25:00
#PBS -l nodes=1:ppn=1
#PBS -M JohnDoe@gmail.com
#PBS -m bea

cd $PBS_O_WORKDIR

### Run gethostname

time ./gethostname.exe
```

To submit the job run:

qsub job.gethostname.bash

PBS Script for MPI Program

```
#!/bin/bash
#PBS -N gethostmpi
#PBS -l nodes=2:ppn=8,pmem=1gb
#PBS -S /bin/bash
#PBS -l walltime=00:05:00
#PBS -M monkeybrain@sc.edu
#PBS -m bea

cd $PBS_O_WORKDIR
##set up your environment
source /etc/profile.d/modules.sh

module add openmpi

mpirun –np 16 ./gethostname.mpi.exe > output.from.txt
```

See job filename: job.gethostname.mpi.bash

To submit the job run:

qsub job.gethostname.mpi.bash

PBS Script for GPGPU Program

Example GPU Job Script

```
#!/bin/bash
#PBS -q gpu
#PBS -N sortrandomkeys
#PBS -l walltime=00:05:00,nodes=1:ppn=4
#PBS -M monkeybrain@gmail.com
#PBS -m bea

cd $PBS_O_WORKDIR
##set up your environment
source /etc/profile.d/modules.sh
module add cuda-toolkit
time ./sortkeys_basic_thrust.exe
```

Give me a node with
this kind of gpu

Load CUDA
environment

See job filename: job.gpu.bash or job.gpu.sort.bash

PBS Script for Executing Matlab script

Using Proprietary Matlab runtime (sample I)

```
#!/bin/bash
#PBS -N matlabjob
#PBS -l nodes=1:ppn=1,pmem=1gb
#PBS -S /bin/bash
#PBS -l walltime=00:05:00

cd $PBS_O_WORKDIR

##set up your environment
source /etc/profile.d/modules.sh
module add matlab

matlab << EOF
a = 10; b = 20; c = 30;
d = sqrt((a + b + c)/pi)
exit
EOF
```

See job filename: job.matlab1.bash

PBS Script for Executing Matlab script

Using Proprietary Matlab runtime (sample II)

```
#!/bin/bash
#PBS -N matlabjob
#PBS -S /bin/bash
#PBS -I walltime=20:05:00,pmem=1gb,nodes=1:ppn=4
#PBS -M monkeybrain@uh.edu
#PBS -m bea
##set up your environment
cd $PBS_O_WORKDIR
source /etc/profile.d/modules.sh
module add matlab

## run matlab program compute pseudo inverse for 100 matrices of size 400x400
time matlab < matrix_inversion100_matlab.m
```

See job filename: job.matlab2.bash

PBS Script for Executing R script

```
#!/bin/bash
#PBS -N R-job
#PBS -S /bin/bash
#PBS -I walltime=0:05:00,pmem=1gb,nodes=1:ppn=1
#PBS -M monkeybrain@uh.edu
#PBS -m bea
##set up your environment
cd $PBS_O_WORKDIR
source /etc/profile.d/modules.sh
module add R

## run R program
R -vanilla < sample.r
```

See job filename: job.R.bash

PBS Script for NAMD2 Job Parallel MPI

```
#!/bin/bash
#PBS -l nodes=1:ppn=4
#PBS -l walltime=00:05:00
#PBS -l pmem=1gb
#PBS -N namd
#PBS -V

cd $PBS_O_WORKDIR

module load namd

mpirun -v namd2 apoa1.namd
```

See job filename: job.namd.mpi.bash

Killing Jobs

- ☞ One, or a few jobs:
 - `qdel [jobID] [jobID] [jobID] ...`
- ☞ Kill all of your jobs:
 - `qselect -u $USER | xargs qdel`
- ☞ Kill all of your queued jobs:
 - `qselect -u $USER -s Q | xargs qdel`
- ☞ Kill all of your running jobs:
 - `qselect -u $USER -s R | xargs qdel`

Try An Example Job

- ☞ Several job scripts files available to use for practice

- ☞ Edit job.**.bash files as needed, then run the job:
 - you can use gedit

- ☞ submit the job(s) i.e.
 - qsub job.**.bash
 - where ** represents the job pattern you are going to submit

job.gethostname.bash job.gethostname.mpi.bash job.gpu.bash job.matlab.bash job.namd.bash

Homework

- ❖ Examine the job files and fix potential glitches
- ❖ Submit the jobs.
- ❖ Record the runtime or execution time and send me these timings as well as the job log files

Maxwell User Support

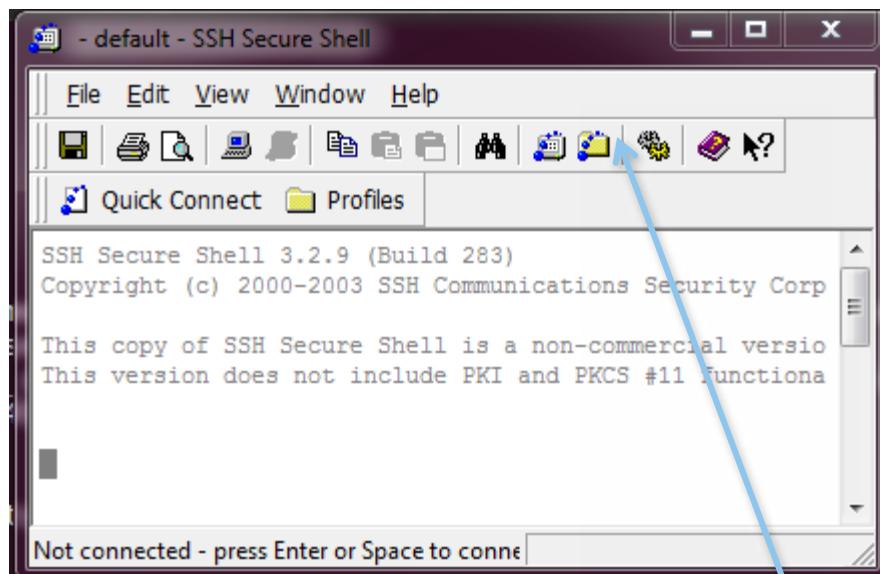
- ☞ Computational scientists are available for
 - consultation and training
 - We can help with porting your homebrew code and domain applications to run on the cluster
 - We can also help with developing proposals that involve use of HPC (i.e. make use of Maxwell resources).
- ☞ “Help” requests or technical questions can be submitted to the Maxwell admin staff and anyone who can help will respond.
- ☞ Visit → <http://support.cacds.uh.edu>
 - To Open A Support Ticket

Acknowledgements

- » Center for Advanced Computation and Data Systems
 - Tony Curtis
 - Barbara Chapman

Accessing Maxwell from Windows

- ❖ Command-line interface, any Secure Shell (ssh) client i.e. putty



- ❖ Transfer files to/from using FileZilla (or Winscp)

Getting Started

- ☞ Open up a gnome terminal

- Multiple ways to do this:
 - Click on the gnome-terminal icon on the top panel
- Or
 - On mouse right click and select terminal and click

User name = student0XX

where XX = 1 → 48

Password = provided

☞ ssh [username]@208.117.189.1

The screenshot shows a terminal window titled "jebalunode@c01:~". Inside the terminal, the user is performing an SSH connection to another host. The command entered is "ssh student5@planck.psc.sc.edu". The terminal displays a warning about host authenticity, asking if the user wants to continue connecting. The user responds "yes", and the terminal adds the host to the list of known hosts. It then prompts for the password of the user "student5@planck.psc.sc.edu". Finally, it displays a welcome message from the cluster: "Welcome to the planck cluster". At the bottom of the terminal, there is a note: "Please review <http://imi.cas.sc.edu/acm-chem/wiki/ClusterHardware> to see the specs on the planck cluster".

```
[jebalunode@c01 ~]$ ssh student5@planck.psc.sc.edu
The authenticity of host 'planck.psc.sc.edu (129.252.37.217)' can't be established.
RSA key fingerprint is 1c:08:3b:30:2f:10:2d:5b:09:f0:ec:52:a0:e6:ce:8b.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'planck.psc.sc.edu,129.252.37.217' (RSA) to the list
of known hosts.
student5@planck.psc.sc.edu's password:
Welcome to the planck cluster

Please review http://imi.cas.sc.edu/acm-chem/wiki/ClusterHardware
to see the specs on the planck cluster
[student5@c01 ~]$
```

Serial Task Management

Use `pbsdsh` to distribute groups of serial tasks within the PBS environment.

Within your “master” script (the one you submit with `qsub`):

```
pbsdsh -v $PBS_O_WORKDIR/subscript.bash
```

This runs `subscript.bash` on each core in your allocation:

```
#!/bin/bash
cd $PBS_O_WORKDIR
myexe < ${PBS_VNODENUM}.inp > ${PBS_VNODENUM}.out
```