

# Introduction to Numerical Libraries

**Jerry Ebalunode**

Center for Advanced Computation and Data Systems  
(CACDS)

<http://cacds.uh.edu>

<http://support.cacds.uh.edu>

University of Houston



@UHCACDS 

[facebook.com/CACDSATUH](https://facebook.com/CACDSATUH) 

# Overview

- Scientific Computing
- Why Numerical Libraries
- Random Number Generation (Intel MKL)
- Vector Arithmetic ( Intel MKL)
- Fast Fourier Transforms ( FFTW & MKL)
- Linear Algebra ( BLAS and LAPACK)
- Open Lab and Homework

# First Access Your Account

- Log into your accounts
  - Username or login = hpc\_user**X**
  - Where **x** = sign in serial number 1 – 47
  - Password = **cacds2014**
  - Use your web browser
    - Firefox, Chromium or Google chrome
- Slides could be downloaded from URL below
  - <http://129.7.249.171/workshops/intro2numlib.pdf>

# Getting Started

Use the terminal to download intro2numlib.zip file to your home directory

- Run the following commands

```
cd
```

```
wget http://129.7.249.171/workshops/intro2numlib.zip
```

```
unzip intro2numlib.zip
```

```
cd intro2numlib
```

Now, you can begin working with tutorial files on your terminal

# THE GRAND CHALLENGE EQUATIONS

$$B_i A_i = E_i A_i + \rho_i \sum_j B_j A_j F_{ji} \quad \nabla \times \vec{E} = - \frac{\partial \vec{B}}{\partial t} \quad \vec{F} = m \vec{a} + \frac{dm}{dt} \vec{v}$$

$$dU = \left( \frac{\partial U}{\partial S} \right)_V dS + \left( \frac{\partial U}{\partial V} \right)_S dV \quad \nabla \cdot \vec{D} = \rho \quad Z = \sum_j g_j e^{-E_j/kT}$$

$$F_j = \sum_{k=0}^{N-1} f_k e^{2\pi i j k / N} \quad \nabla^2 u = \frac{\partial u}{\partial t} \quad \nabla \times \vec{H} = \frac{\partial \vec{D}}{\partial t} + \vec{J}$$

$$p_{n+1} = r p_n (1 - p_n) \quad \nabla \cdot \vec{B} = 0 \quad P(t) = \frac{\sum_i W_i B_i(t) P_i}{\sum_i W_i B_i(t)}$$

$$-\frac{\hbar^2}{8\pi^2 m} \nabla^2 \Psi(r, t) + V \Psi(r, t) = -\frac{\hbar}{2\pi i} \frac{\partial \Psi(r, t)}{\partial t} \quad -\nabla^2 u + \lambda u = f$$

$$\frac{\partial \vec{u}}{\partial t} + (\vec{u} \cdot \nabla) \vec{u} = -\frac{1}{\rho} \nabla p + \gamma \nabla^2 \vec{u} + \frac{1}{\rho} \vec{F} \quad \frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} + \frac{\partial^2 u}{\partial z^2} = f$$

• NEWTON'S EQUATIONS • SCHROEDINGER EQUATION (TIME DEPENDENT) • NAVIER-STOKES EQUATION •  
 • POISSON EQUATION • HEAT EQUATION • HELMHOLTZ EQUATION • DISCRETE FOURIER TRANSFORM •  
 • MAXWELL'S EQUATIONS • PARTITION FUNCTION • POPULATION DYNAMICS •  
 • COMBINED 1ST AND 2ND LAWS OF THERMODYNAMICS • RADIOSITY • RATIONAL B-SPLINE •

**SAN DIEGO SUPERCOMPUTER CENTER**

*A National Laboratory for Computational Science and Engineering*

[Courtesy of San Diego Supercomputer Center]

# Scientific Computing

- Why should we care about scientific computing?
  - Computational research has emerged to complement experimental methods in basic research, design, optimization, and discovery in all facets of engineering and science
  - In certain cases, computational simulations are the only possible approach to analyze a problem:
    - Experiments may be cost prohibitive (e.g. *flight testing a 1,000 fuselage/wing-body configurations for a modern fighter aircraft*)
    - Experiments may be impossible (e.g. *interaction effects between the International Space Station and Shuttle during docking*)
  - Simulation capabilities rely heavily on the underlying compute power (e.g. amount of memory, total compute processors, and processor performance)
    - Fostered the introduction and development of *super-computers* starting in the 1960's
    - Large-scale compute power is tracked around the world via the *Top500 List* (more on that later)

[Courtesy of San Diego Supercomputer Center]

# Scientific Computing: a definition

- “The efficient computation of constructive methods in applied mathematics”
  - Applied math: getting results out of application areas
  - Numerical analysis: results need to be correctly and efficiently computable
  - Computing: the algorithms need to be implemented on modern hardware

[Courtesy of San Diego Supercomputer Center]

# Case Study

## Random Number Generation

- A random number generator (RNG) is a computational device designed to generate a sequence of numbers or symbols that lack any pattern, i.e. appear random.
- RNGs are widely used in Monte Carlo-method simulations, molecular dynamics simulations, cryptography etc.



# Random Number Generation using GNU C library

```
#include <iostream>
#include <omp.h>
using namespace std;
int main()
{
    double t1,t0,elapsed;
    const size_t N = 1<<29L;
    //const size_t F = sizeof(float);
    float* A = new float [N];
    srand(0); // Initialize RNG
    t0=omp_get_wtime();
    for (int i = 0; i < N; i++)
    {
        A[i]=(float)rand() / (float)RAND_MAX;
    }
    t1=omp_get_wtime();
    elapsed=t1-t0;

    cout << "\nGenerated " << N << " random numbers in " << elapsed << " seconds\n\nHere is a sample\n";

    for (int i = 0; i < 10; i++)
    {
        cout << endl << A[i];
    }
    delete [] A;
}
```

# Random Number Generation

```
module add intel
```

```
icpc random_gen.cpp -openmp -o random_gen
```

```
./random_gen
```

Generated 536,870,912 random numbers in 5.49038 seconds

Here is a sample

0.242578

# Numerical Libraries

- Software libraries used in application development for performing numerical calculations.
  - Robust and efficient algorithms
- Benefits
  - Helps scientists from reinventing the wheel
  - Scientist can focus more on the original problem
- Example
  - OpenSource
    - Blas, Lapack, GSL, GMP, FFTW, ACML, SuitSparse, Magma
  - Commercial
    - Intel Math Kernel Library, NAG, IMSL, ALGLIB

# Intel Math Kernel Library

- MKL improves performance with math routines for software applications that solve large computational problems.
- MKL provides:
  - BLAS and LAPACK linear algebra routines
  - Fast Fourier transforms
  - Vectorized math functions
  - Random number generation functions
  - Tools for solving partial differential equations

# Random Number Generation using Intel Math Kernel Library

A typical algorithm for MKL random number generators is as follows:

- Create and initialize stream/streams. Functions `vsNewStream`
- Call one or more RNGs.
- Process the output.
- Delete the stream/streams. Function `vsDeleteStream`.
- Reference:
  - <https://software.intel.com/en-us/node/521842>

# Random Number Generation using Intel Math Kernel Library

```
#include <iostream>
#include <mkl_vsl.h>
#include <omp.h>
using namespace std;
int main()
{
    double t1,t0,elapsed;
    const size_t N = 1<<29L;
    //const size_t F = sizeof(float);
    float* A = new float [N];

    VSLStreamStatePtr rnStream;
    vslNewStream(&rnStream, VSL_BRNG_MT19937, 1 );      ; // Initialize RNG

    t0=omp_get_wtime();
    vsRngUniform(VSL_RNG_METHOD_UNIFORM_STD, rnStream, N, A, 0.0f, 1.0f);
    t1=omp_get_wtime();
    elapsed=t1-t0;

    cout << "\nGenerated " << N << " random numbers in " << elapsed << " seconds\n\nHere is a sample\n";
    for (int i = 0; i < 10; i++)
    {
        cout << endl << A[i];
    }
    delete [] A;
    vslDeleteStream( &rnstream );
}
```

# Random Number Generation

```
icc -openmp -mkl=sequential random_gen_numlib.cpp -o random_gen_numlib
```

```
./random_gen_numlib
```

Generated 536,870,912 random numbers in  
**1.57736** seconds

~4X speedup

Here is a sample

0.134364

# Vector Arithmetic

```
for (int i = 0; i < N; i++)  
{  
    C[i] = A[i] + B[i];  
}
```



# Vector Add

```
icpc vecadd.cpp -openmp -o vecadd
```

```
./vecadd
```

# Vector Arithmetic with Intel MKL

`vsAdd( N, A, B, C)`

# Vector Add

```
icpc  vecadd_numlib.cpp -openmp -o  
vecadd_numlib -mkl=sequential
```

```
./vecadd_numlib
```

# Fast Fourier Transform

- FFT is an algorithm to compute Discrete Fourier Transforms (DFT) in a fast way
- FFT employed in simulation of periodic systems
  - periodic systems are very common in scientific computing
    - many body systems in empirical and *ab initio* molecular dynamics
- DFT is simply the computation of the coefficients  **$c_k$** , the integrals, using the trapezoidal integration formula, that is, if  $x_0, x_1, \dots, x_{N-1}$  are  $N$  complex numbers that represents  $f(n/N) = x_n$ , then

$$\sum_{k=-\infty}^{\infty} c_k e^{2\pi k i \theta} : c_k = \int_{\mathbb{T}} e^{-2\pi k i \theta} f(\theta) d\theta : c_k = \frac{1}{N} \sum_{n=0}^{N-1} x_n e^{-ik \frac{2\pi n}{N}}$$

# Fast Fourier Transform

## 1D DFT

```
int main(int argc, char* argv[])
{
    double t1,t0,elapsed;
    const size_t N = 1<<25L;
    fftw_complex *in, *out;
    fftw_plan p;
    int i;
    float *real = new float [N];
    float *imag = new float [N];
    VSLStreamStatePtr rnStream;
    vslNewStream(&rnStream, VSL_BRNG_MT19937, 1 );      ; // Initialize RNG
    in = (fftw_complex*) fftw_malloc(sizeof(fftw_complex)*N);
    vsRngUniform(VSL_RNG_METHOD_UNIFORM_STD, rnStream, N, real, 0.0f, 1.0f);
    vsRngUniform(VSL_RNG_METHOD_UNIFORM_STD, rnStream, N, imag, 0.0f, 1.0f);
    for(i=0;i<N;i++)
    {
        in[i][0]=real[i];
        in[i][1]=imag[i];
    }
    out = (fftw_complex*) fftw_malloc(sizeof(fftw_complex)*N);
    //! STEP ONE
    p = fftw_plan_dft_1d(N,in,out,FFTW_FORWARD,FFTW_ESTIMATE);
```

# Fast Fourier Transform

```
//! STEP TWO
t0=omp_get_wtime();
fftw_execute(p);
t1=omp_get_wtime();
elapsed=t1-t0;

printf("1-D FFT is:\n");
for(i=0;i<10;i++)
{
    cout << out[i][0]<<"\t"<< out[i][1] <<endl;
}
//! STEP THREE
fftw_destroy_plan(p);
fftw_free(in);
fftw_free(out);
delete [] real;
delete [] imag;
vslDeleteStream(&rnStream);
return 0;
}
```

# FFT example

```
icpc fftw_example.cpp -openmp -o  
fftw_example -mkl=sequential
```

```
./fftw_example
```

# BLAS

## Basic Linear Algebra Subprograms

- The Basic Linear Algebra Subprograms (BLAS) define a set of fundamental operations on vectors and matrices which can be used to create optimized higher-level linear algebra functionality.
- Tuned Linear Algebra Programs
- **Level 1**
  - Vector operations, e.g.  $y = \alpha x + y$
- **Level 2**
  - Matrix-vector operations, e.g.  $y = \alpha A x + \beta y$
- **Level 3**
  - Matrix-matrix operations, e.g.  $C = \alpha A B + C$



# BLAS

The BLAS routines and functions are divided into the following groups according to the operations they perform:

- BLAS Level 1 Routines perform operations of both addition and reduction on vectors of data. Typical operations include scaling and dot products.
- BLAS Level 2 Routines perform matrix-vector operations, such as matrix-vector multiplication, rank-1 and rank-2 matrix updates, and solution of triangular systems.
- BLAS Level 3 Routines perform matrix-matrix operations, such as matrix-matrix multiplication, rank-k update, and solution of triangular systems.

# BLAS

## SAXPY

```
#include "mkl_cblas.h"

#include <omp.h>
using namespace std;

int main(int argc, char* argv[])
{
    double t1,t0,elapsed;
    float alpha=3.0f;
    MKL_INT i, N = 1<<25L;
    //! STEP ZERO
    float *x = new float [N];
    float *y = new float [N];
    VSLStreamStatePtr rnStream;
    vslNewStream(&rnStream, VSL_BRNG_MT19937, 1 );      ; // Initialize RNG

    vsRngUniform(VSL_RNG_METHOD_UNIFORM_STD, rnStream, N, x, 0.0f, 1.0f);
    vsRngUniform(VSL_RNG_METHOD_UNIFORM_STD, rnStream, N, y, 0.0f, 1.0f);

    //! STEP ONE
    t0=omp_get_wtime();
    cblas_saxpy(N, alpha, x, 1, y, 1);
    t1=omp_get_wtime();
```

BLAS CALL



# BLAS

## SAXPY

```
t1=omp_get_wtime();
elapsed=t1-t0;

printf("SAXPY is:\n");
for(i=0;i<10;i++)
{
    cout << y[i] <<endl;
}
//! STEP TWO
delete [] x;
delete [] y;
vslDeleteStream(&rnStream);
return 0;
}
```

# SAXPY example

```
icpc saxpy.cpp -openmp -o saxpy -  
mkl=sequential
```

```
./saxpy
```

# LAPACK

- The original versions of LAPACK from which that part of Intel MKL was derived can be obtained from <http://www.netlib.org/lapack/index.html>. The authors of LAPACK are E. Anderson, Z. Bai, C. Bischof, S. Blackford, J. Demmel, J. Dongarra, J. Du Croz, A. Greenbaum, S. Hammarling, A. McKenney, and D. Sorensen.
- The LAPACK routines can be divided into the following groups according to the operations they perform:
  - Routines for solving systems of linear equations, factoring and inverting matrices, and estimating condition numbers (see LAPACK Routines: Linear Equations).
  - Routines for solving least squares problems, eigenvalue and singular value problems, and Sylvester's equations (see LAPACK Routines: Least Squares and Eigenvalue Problems).
  - Auxiliary and utility routines used to perform certain subtasks, common low-level computation or related tasks (see LAPACK auxiliary and utility routines).

# LAPACK Routines

- The library includes LAPACK routines for both real and complex data. Routines are supported for systems of equations with the following types of matrices:
  - general
  - banded
  - symmetric or Hermitian positive-definite (full, packed, and rectangular full packed (RFP) storage)
  - symmetric or Hermitian positive-definite banded
  - symmetric or Hermitian indefinite (both full and packed storage)
  - symmetric or Hermitian indefinite banded
  - triangular (full, packed, and RFP storage)
  - triangular banded
  - tridiagonal
  - diagonally dominant tridiagonal.

# Routine Naming Conventions

- To call one of the routines from a FORTRAN 77 program, you can use the LAPACK name.
  - LAPACK names have the structure **?yyzzz** or **?yyzz**, where the initial symbol ? indicates the data type:
    - s real, single precision
    - c complex, single precision
    - d real, double precision
    - z complex, double precision

# Eigen value and Eigen vectors using LAPACK Drivers

- **ssyev**: computes eigenvalues and, optionally, the eigenvectors of a square real symmetric matrix  $A$
- *interface*
- *ssyev(jobz, uplo, n, a, lda, w, work, lwork, info)*



# Eigen value and Eigen vectors using LAPACK Drivers

```
ssyev( "Vectors", "Upper", &n, a, &lda, w,  
&wkopt, &lwork, &info );
```

```
lwork = (int)wkopt;
```

```
work = (float*)malloc( lwork*sizeof(float) );
```

```
/* Solve eigenproblem */
```

```
ssyev( "Vectors", "Upper", &n, a, &lda, w, work,  
&lwork, &info );
```

# SSYEV example

```
icpc ssyev.cpp -o ssyev -mkl=sequential
```

```
./ssyev
```

# Parallel Numerical Library

- Three flavors
  - Multi-threaded support for multicore platforms
    - OpenMP for shared memory systems
  - Distributed memory support
    - MPI enabled libraries
      - SCALAPACK
- GPGPU i.e. support Accelerators and Co-Processors
  - Xeon Phi Intel MKL offload support
  - Magma library <http://icl.cs.utk.edu/magma/>
    - BLAS and LAPACK support on NVIDIA and Intel GPGPUs
    - open source
  - CuBLAS, CuFFT, CuSparse
    - Support NVIDIA GPUs

# Parallel Numerical Library

## Multi-core Platforms

Multi-threaded support for multicore platforms

- OpenMP for shared memory systems
- Example with ssyev

```
icpc ssyev.cpp -o ssyev.parallel -mkl=parallel
```

```
export OMP_NUM_THREADS=#CORES_ON_SERVER  
./ssyev
```

# HPC Support Page

- <http://support.cacds.uh.edu/>

