

# Introduction to Cluster Computing @ UH Using Opuntia and Maxwell

**Jerry Ebalunode**  
[jebalunode@uh.edu](mailto:jebalunode@uh.edu)

**Center for Advanced Computation and Data Systems  
(CACDS)**  
<http://cacds.uh.edu>



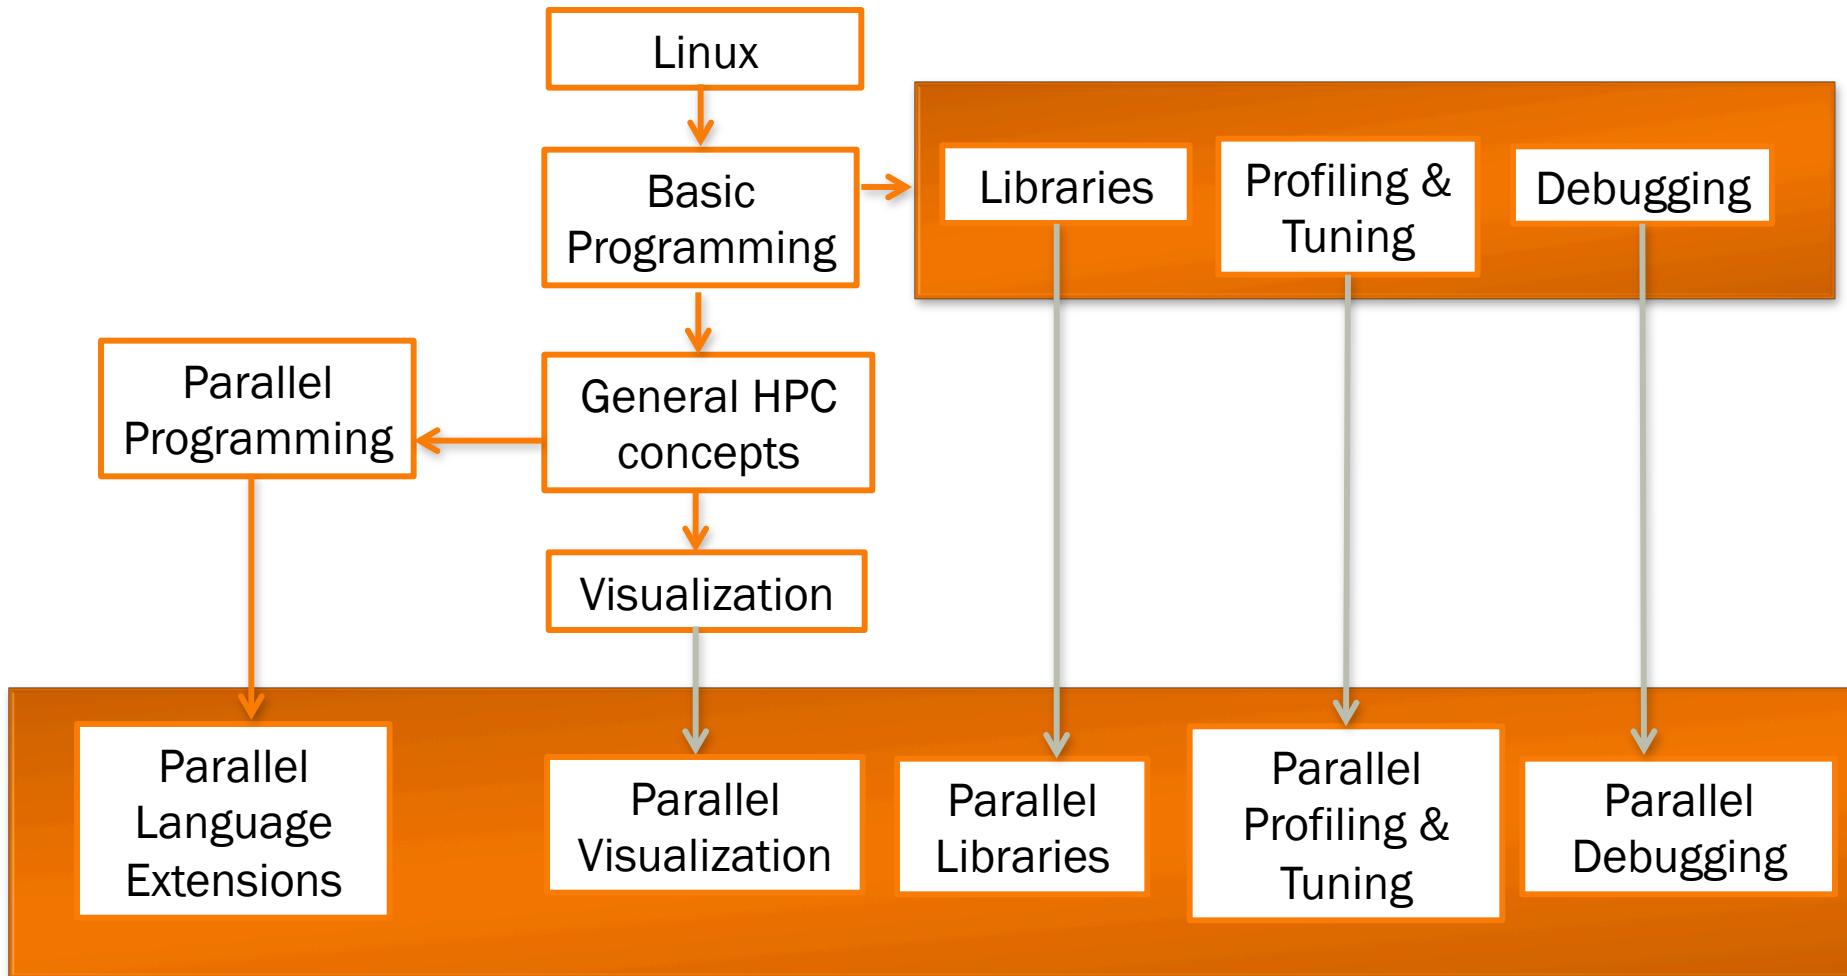
**University of Houston**  
Houston, TX

# CACDS

## Mission Statement

- ∞ The Center for Advanced Computing and Data Systems (CACDS) provides high performance computing resources to advance Tier One research and education goals at the University of Houston (UH).

# HPC User Training Road Map



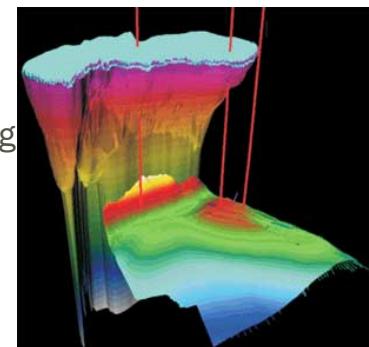
# HPC Training

**Location: PGH 200**  
**48 seat class room**

- Linux
- Shell programming
- Cluster computing on Maxwell
- XSEDE Cyberinfrastructure and allocation process
- Python
- MATLAB
- C++
- Fortran
- Parallel Programming using OpenMP
- Parallel Programming using MPI
- Parallel numerical libraries

**Location: PGH 235**  
**12 seat class room**

- Nvidia CUDA
  - introductory, intermediate and advanced
  - profiling, debugging & tuning
- OPENACC
  - intermediate, advanced
  - profiling, debugging & tuning
- Intermediate and Advanced MPI
  - profiling, debugging & tuning
- Intel Xeon Phi Programming
  - Native computing
    - OpenMP, Cilk++
  - Offload Execution
  - Symmetric Execution
  - profiling, debugging & tuning
- Visualization at scale
  - Paraview
  - VisIt



# CACDS Training

## [www.cacds.uh.edu/training](http://www.cacds.uh.edu/training)

UNIVERSITY **of HOUSTON**  
CENTER FOR ADVANCED COMPUTING & DATA SYSTEMS

ABOUT ▾ RESEARCH ▾ EDUCATION ▾ RESOURCES ▾ NEWS ▾



## | Training Courses

**EDUCATION**

Training Courses>>

- [Events](#)
- [Lectures](#)
- [Workshops](#)
- [Technical Seminars](#)

CACDS has launched several HPC training courses for members of the UH research community. Registration is required to take the courses. Please click the date you would like to attend under each course to register.

**CURRENT TRAINING**

**CLASS NAME**

[Introduction to Linux](#)

[Introduction to Shell Programming](#)

[Introduction to Cluster Computing and Maxwell Cluster](#)

[Introduction to C++ Programming](#)

[Introduction to Python Programming](#)

[Introduction to Parallel Programming with OpenMP](#)

[GPGPU Parallel Programming with OpenACC](#)

[GPGPU Parallel Programming with CUDA](#)

[Introduction to Parallel Programming with MPI](#)

[Introduction to Matlab](#)

[Introduction to Intel Xeon Phi Programming](#)

[Introduction to High Performance Numerical Libraries](#)

[Visualization with VISIT I](#)

# Overview

## Opuntia Cluster

- Accessing the Cluster
- Maxwell Specifications
- User Environment
- Compiling Programs
- Portable Batch System (PBS)
- Batch Jobs
- Interactive Jobs
- Killing & Suspending Jobs

## Maxwell Cluster

- Accessing the Cluster
- Maxwell Specifications
- User Environment
- Compiling Programs
- Portable Batch System (PBS)
- Batch Jobs
- Interactive Jobs
- Killing & Suspending Jobs

# First Access Your Account

- ☞ Log into your accounts

- Username or login = hpc\_userX
- Where x = sign in serial number 1 – 47
- Password = **cacds2014**
- Use your web browser
  - Firefox, Chromium or Google chrome

- ☞ Slides could be downloaded from URL below

**/share/apps/tutorials/intro2Opuntia.pdf**

# Disclaimer

- ☞ Maxwell & Opuntia clusters are for research work under a UH PI
  - Faculty sponsored RESEARCH activity only
- ☞ You will be using a surrogate cluster for this training
  - You would not be using Maxwell for this training
- ☞ We are assuming you are already familiar with the Linux Operating Environment
  - If not try to attend any of the upcoming Linux hands on tutorials offered by CACDS
- ☞ After this training, if you do not already have an account on Maxwell/Opuntia feel free to request one, after authorization from your supervisor (UH Faculty)
  - <http://www.rcc.uh.edu/cgi-bin/account.cgi> (actual application page)
  - <http://www.cacds.uh.edu/OpuntiaAccountRequest>

# Who Can Use the Clusters?

- ∞ Open to all UH Central faculty for research purposes
  - Students & postdocs – under UH faculty guidance
  - Collaborators – under UH faculty guidance

# What Happens Next?

## ☞ How to connect to the Opuntia Cluster:

- For UNIX/Linux/Mac users
  - Open up terminal console and use Openssh secure shell client to connect
  - syntax:
    - **ssh username@opuntia.cacds.uh.edu**
    - **For X11 support**
      - **ssh -YX username@opuntia.cacds.uh.edu**

Where **username** is your assigned login ID..

- you might be prompted to accept finger print...
  - Just type “yes” followed by Enter key
- For Windows users
  - Install Putty or any other compatible SSH client to connect
  - Putty URL = <http://www.putty.org/>

# Connecting via OpenSSH Client on UNIX based Systems

```
129-7-249-201:~ jebalunode$ ssh -Y opuntia.cacds.uh.edu -l jebaluno
The authenticity of host 'opuntia.cacds.uh.edu (129.7.54.49)' can't be established.
RSA key fingerprint is a7:e7:7f:16:04:cd:90:23:b7:3a:a0:05:00:c4:41:fc.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'opuntia.cacds.uh.edu,129.7.54.49' (RSA) to the list of known hosts.
```

Use of University of Houston computing and network facilities requires prior authorization. Unauthorized access is prohibited. Usage may be subject to security testing and monitoring. Abuse is subject to criminal prosecution. A complete manual of security policies and procedures is available at <http://www.uh.edu/> in the Administration directory.

Shared  
Password:

# Connecting via OpenSSH Client on UNIX based Systems

**Password:**

Last login: Mon Feb 2 17:17:01 2015 from c-73-55-246-235.hsd1.tx.comcast.net  
Rocks 6.1.1 (Sand\_Boa)  
Profile built 11:03 26-Sep-2014

Kickstarted 06:43 26-Sep-2014

Turn Paper into Digital  
Scan paper documents into PDF  
Easily and quickly and accurately.

[Learn How](#) Opuntia Cluster

<http://en.wikipedia.org/wiki/Opuntia>

Shared

This cluster uses SLURM as its job scheduler

Use module avail to list available software

Use module add <name> to set up environment to use <name>

Example module add lammps

Green

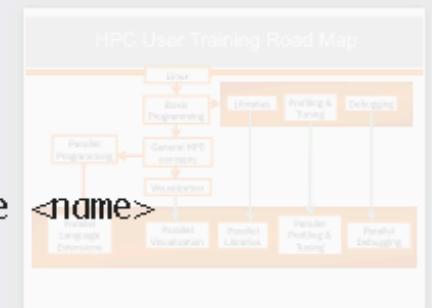
Blue

Purple

Gray

All Tags...

[jebaluno@opuntia ~]\$ █



# Accessing Tutorial Materials

First login into the cluster:

**TYPE AND EXECUTE COMMANDS IN RED!!!**

ssh -Y 10.1.1.1

cd

cp /share/apps/tutorials/intro2maxwell.zip ~

unzip intro2maxwell.zip

cd intro2maxwell

# The Opuntia Cluster

- ∞ Both Shared and Distributed systems
  - single core, multicore (multithreaded) and MPI jobs
- ∞ 80+ compute nodes (1700 + cpu cores)
  - <http://baragon.hpcc.uh.edu/ganglia/?c=Opuntia>
  - 2 Xeon Phi 5110P
  - 4 Nvidia Tesla K40 GPUs
- ∞ Operating system: Linux
  - Distribution: Redhat 6.5
- ∞ 56 Gigabit network interconnect
- ∞ Data Storage:
  - Default: 10+ GB Home directory space per user
  - PI's could request for project directory space to host files needed for simulations
    - Justification might be required

# Opuntia Specifications

**Hardware Specification Table**

Node Type	CPU Type	CPU Socket Count	Total Cores	Memory	Disk Space	Node Count
Login HP DL380	Intel Xeon E5-2680v2 2.8 GHz	2	20	64 GB	2.4 TB	1
Compute HP SL230	Intel Xeon E5-2680v2 2.8 GHz	2	20	64 GB	1 TB	80
Large Memory HP DL560	Intel Xeon E5-4650v2 2.4 GHz	4	40	512 GB	1 TB	2
XLarge Memory HP DL580	Intel Xeon E7-4880v2 2.5 GHz	4	60	1 TB	1 TB	1
GPU Accelerator HP SL250	Intel Xeon E5-2680v2 2.8 GHz	CPU:2	CPU:20	CPU:64 GB	1TB	2
	Tesla K40 GPU	GPU:2	GPU:5760	GPU:24 GB		
Xeon Phi Coprocessor HP SL250	Intel Xeon E5-2680v2 2.8 GHz	CPU:2	CPU:20	CPU:64 GB	1TB	2
	Xeon Phi 5110P	Phi:1	Phi:61	Phi:8 GB		
Storage HP SL4540	Intel Xeon E5-4650v2 2.4 GHz	2	20	64 GB	120 TB	4

## Storage

384 TB of NFS storage

## Interconnect

Opuntia nodes are connected via 56 Gb/s Ethernet interconnect.

# HPC Software Available on Opunta

## ∞ Proprietary & Open-Source Software:

- Examples:
  - Amber, Autodock, Bwa, NwChem, Espresso, Octave, Matlab, Gromacs, Lammmps, NAMD , R, NAG, MKL , ASE, GSL, OpenMPI, Intel, Gnu and PGI compilers ....
  - and much more...
- Proprietary software may be limited to licensed users e.g. VASP
- Open-source software is accessible to all

## ∞ PI could request for installation of additional software

- PI's licensed propriety application can be installed too

## ∞ Software environment management via “**modules**”

- Allows software to be accessed via an assigned module
- Dynamic modification of a user's environment

# Computing Environment Setup Using Modules

- ❖ The module system is used to make software and related settings available easily
  - software environments can be loaded and unloaded dynamically

- ❖ To get a list of available software:

**module avail**  
(also useful: **module list**)

- ❖ To clear-out all added modules:

**module purge**

- ❖ To clear-out all specific modules:

**module rm module1 module2 ....**

- ❖ To use specific software every time you log-in, place the

module add command in your `~/.bashrc` file i.e:

**echo "module add intel intelmpi" >> ~/.bashrc**

- ❖ Now add intel compiler and MPI runtime (Openmpi version) to your environment)

**module add intelmpi**

# Compiling Programs

- ❖ CPU programs (compiled to execute on central processing unit)
  - Serial Programs
  - Parallel Programs
- ❖ GPU programs (compiled to execute on graphics processing unit)
  - GPGPU Programs

# Compiling Programs: CPU Serial Program

## Example 1 Program in file: gethostname.c

```
#include <stdio.h>
#include <sys/utsname.h>
int main ( )
{
    struct utsname uts;
    uname (&uts);
    printf ("Process on node %s.\n",
    uts.nodename);
    return 0;
}
```

- To compile:
  - 1. choose compiler  
`module add intel`
  - 2. Compile  
`icc gethostname.c -o gethostname.exe`

# Compiling Programs: CPU Parallel Program

## Programs using Message Passing Interface (MPI)

Example: 1 Program in file : gethostname.mpi.c

```
#include <stdio.h>
#include <sys/utsname.h>
#include <mpi.h>
int main (int argc, char *argv[]){
    struct utsname uts;
    int rank;
    MPI_Init (&argc, &argv);
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    uname (&uts);
    printf ("Process %d on node %s.\n", rank, uts.nodename);
    MPI_Finalize ();
    return 0;
}
```

- To compile:
  - 1. choose compiler  
`module add openmpi`
  - 2. Compile  
`mpicc gethostname.mpi.c -o gethostname.mpi.exe`

# Compiling Programs: GPGPU Parallel Program

## Example 1 Program in file: sortkeys\_basic\_thrust.cu

```
#include <thrust/host_vector.h>
#include <thrust/device_vector.h>
#include <thrust/sort.h>
int main(void)
{
    // generate 16 Million random numbers on the host
    thrust::host_vector<int> h_vec(1 << 24);
    thrust::generate(h_vec.begin(), h_vec.end(), rand);
    thrust::device_vector<int> d_vec= h_vec; // transfer data to the device
    thrust::sort(d_vec.begin(), d_vec.end()); // sort data on the device (805 Mkeys/sec on GeForce GTX 480)
    thrust::copy(d_vec.begin(), d_vec.end(), h_vec.begin()); // transfer data back to host
    return 0;
}
```

- To compile:
  - 1. choose compiler

```
module add cuda-toolkit
```

```
2. Compile name = nvcc
```

```
nvcc sortkeys_basic_thrust.cu -o sortkeys_basic_thrust.exe
```

# The Queue System (SLURM)



- ❖ Simple Linux Utility for Resource Management (aka SLURM) is the resource management service used on the Opuntia Cluster
- ❖ Enables you to make more efficient use of your time through scripting computational tasks
- ❖ SLURM takes care of running these tasks and returning the results
- ❖ If the cluster is full, SLURM holds your tasks and runs them when the resources are available
- ❖ SLURM ensures fair sharing of cluster resources (policy enforcement)
  - SLURM ensures optimal/efficient use of available resources

# SLURM Commands

Command(s)	Description
<b>squeue</b>	Check status of all jobs
<b>sbatch ./myjob</b>	Submit batch jobs
<b>squeue -u \$USER</b>	Check status of just your jobs
<b>srun --pty /bin/bash -l</b>	Submit an interactive job
<b>srun --x11=first --pty /bin/bash -l</b>	Submit an interactive job <b>with X11 support</b>
<b>scontrol hold jobID</b> i.e. qhold 12345	Put a job on hold (before it starts)
<b>scontrol release jobID</b>	Release a job from hold status
<b>scancel #jobID</b>	Delete a job, running or not



# Using sbatch with a Job Script

- ❖ Syntax/Format:

- **sbatch [sbatch params] [job script file]**

- ❖ Examples:

- **sbatch myjob**

# SLURM Batch Job Script File

- ❖ Very simple example:
- ❖ Serial job requesting 1 CPU, 2 GB memory, 60 hours of walltime

```
#!/bin/bash
#SBATCH -J jobname
#SBATCH -o jobname.o%j
#SBATCH --mem=2gb
#SBATCH -t 60:00:00
```

./my.program.exe [arguments for my program]

# Using Parallel Environment to Manage CPU Core Availability

- ☞ `#SBATCH -N 1 -n 20` --use only 1 node and use a total of 20 processes, 20 CPUs per node.
- ☞ `#SBATCH -N 2 -n 20` --use only 2 nodes and use a total of 20 processes, 10 CPUs per node.
- ☞ `#SBATCH -N 2 -n 40` --use only 2 nodes and use a total of 40 processes, 20 CPUs per node.
- ☞ `#SBATCH -N 4 -n 80` --use only 4 nodes and use a total of 80 processes, 20 CPUs per node.

# Sbatch Notification Parameters

☞ **#SBATCH --mail-user=[e-mail address]**

- e-mail address can be a list of email addresses
- separated by commas.

```
#SBATCH --mail-type=begin # email me when the job starts
```

```
#SBATCH --mail-type=end   # email me when the job finishes
```

```
#SBATCH --mail-type=all   # email me when the job starts & finishes
```

```
#!/bin/bash
```

```
#SBATCH -o jobname.o%j -J o jobname
```

```
#SBATCH --mail-user=jerry@uh.edu
```

```
#SBATCH --mail-type=begin # email me when the job starts
```

```
#SBATCH --mail-type=end   # email me when the job finishes
```

```
./my.program.exe [arguments for my program]
```

# Handling Output Files

- ❖ When a job is started it takes its job name from the script file that was submitted.
  - The standard output and error output are sent to file named jobname.o987349 (in your job directory )
  
- ❖
- ❖ **#SBATCH -o** jobname.o%j name to be used for the job.
- ❖ **#SBATCH -J** jobname name to be used for the job.

# Using Sbatch with a Job Script

- ☞ Syntax/Format:
  - **sbatch [sbatch params] [script file]**
- ☞ More examples on using parameters at the command line:
  - **sbatch -J test1 myjob.sh**
  - **sbatch -N 2 -n 16 -o test2.out -J test2 myjob.sh**
  - **sbatch -t 02:00:00 --mem=4gb myjob.sh**
  - **sbatch -n 32 --mem=64gb myjob.sh**
  - **sbatch -p gpu -N 1 -n 20 my\_gpu\_job.sh**
- ☞ To see all available sbatch options, just run `sbatch --help`

# Using an Interactive Job

- ❖ Special kind of batch job
- ❖ Useful for debugging applications, short tests, or for computational steering

```
srun --pty /bin/bash -l
```

```
srun -n 8 -t 2:00:00 --pty /bin/bash -l
```

```
srun -N 2 -n 8 -t 2:00:00 --pty /bin/bash -l
```

```
srun -n 8 -p gpu -t 2:00:00 --pty /bin/bash -l
```

# Job Queues

Queue Name	Description
short	Queue to run short jobs ( =< 4 hours)
medium	Queue to run medium jobs ( 1 week allowed)
long	Queue to run very long jobs (2 weeks allowed)
gpu	Queue to run gpu jobs on K 40 equipped nodes (24 hrs allowed)

Note phi and gpu queue are not available/functional on the surrogate cluster

# SLURM Script for Serial Program

## 1 CPU

### Example of a Serial Job Script

filename: job.gethostname.bash.slurm

```
#!/bin/bash
#SBATCH -o gethost.out
#SBATCH -t 00:25:00
#SBATCH --mail-user=JohnDoe@gmail.com
#SBATCH --mail-type=all

### Run gethostname

time ./gethostname.exe
```

To submit the job run:

**sbatch job.gethostname.bash.slurm**

# SLURM Script for MPI Program

```
#!/bin/bash
#SBATCH -n 20
#SBATCH -t 0:05:05 # Runtime in D-HH:MM
#SBATCH --mem=1000 # Memory pool for all cores (see also --mem-per-cpu)
#SBATCH -o hostname.out
#SBATCH -e hostname.err
#SBATCH --mail-type=END # Type of email notification- BEGIN,END,FAIL,ALL
#SBATCH --mail-user=jerry@uh.edu
# Email to which notifications will be sent

##set up your environment

module add intelmpi

mpirun -np 16 ./gethostname.mpi.exe > output.from.txt
```

See job filename: job.gethostname.mpi.bash.slurm

To submit the job run:

**sbatch job.gethostname.mpi.bash.slurm**

# SLURM Script for GPGPU Program

## Example GPU Job Script

```
#!/bin/bash
#SBATCH -n 1
#SBATCH -t 0:30:05 -p gpu
#SBATCH --mem=10gb
#SBATCH -o GpuJob.out
#SBATCH --mail-type=all
#SBATCH --mail-user=jerry@uh.edu
```

Give me a node with  
this kind of gpu

##set up your environment

```
module add cuda-toolkit
time ./sortkeys_basic_thrust.exe
```

Load CUDA  
environment

See job filename: job.gpu.bash or job.gpu.sort.bash.slurm

# SLURM Script for Executing Matlab script Using Proprietary Matlab runtime (sample I)

```
#!/bin/bash
#SBATCH -n 8 -N 1
#SBATCH -t 0:30:05
#SBATCH --mem=10gb
#SBATCH -o matlabjob.out -J matlabjob
#SBATCH --mail-type=all
#SBATCH --mail-user=jerry@uh.edu

##set up your environment

module add matlab

matlab << EOF
a = 10; b = 20; c = 30;
d = sqrt((a + b + c)/pi)
exit
EOF
```

See job filename: job.matlab1.bash.slurm

# SLURM Script for Executing Matlab script Using Proprietary Matlab runtime (sample II)

```
#!/bin/bash
#SBATCH -n 8 -N 1
#SBATCH -t 0:30:05
#SBATCH --mem=10gb
#SBATCH -o matlabjob.out -J matlabjob
#SBATCH --mail-type=all
#SBATCH --mail-user=jerry@uh.edu

module add matlab

## run matlab program compute pseudo inverse for 100 matrices of size 400x400
time matlab < matrix_inversion100_matlab.m
```

See job filename: job.matlab2.bash.slurm

# SLURM Script for Executing R script

```
#!/bin/bash
#SBATCH -n 1
#SBATCH -t 0:30:05
#SBATCH --mem=10gb
#SBATCH -o Rjob.out -J Rjob
#SBATCH --mail-type=all
#SBATCH --mail-user=jerry@uh.edu

module add R

## run R program
R --vanilla < sample.r
```

See job filename: job.R.bash.slurm

# SLURM Script for NAMD2 Job Parallel MPI

```
#!/bin/bash
#SBATCH -n 32 -N 4
#SBATCH -t 0:30:05
#SBATCH --mem=10gb
#SBATCH -o NAMDjob.out -J NAMDjob
#SBATCH --mail-type=all
#SBATCH --mail-user=jerry@uh.edu
```

```
module load namd
```

```
mpirun -v namd2 apoa1.namd
```

See job filename: job.namd mpi.bash.slurm

# Killing Jobs

❖ One, or a few jobs:

- `scancel [jobID] [jobID] [jobID] ...`

❖ Kill all of your jobs:

- `scancel -u $USER`

❖ Kill all of your queued jobs:

- `scancel -u $USER -t PENDING`

❖ Kill all of your running jobs:

- `scancel -u $USER -t RUNNING`

# Try An Example Job

- ❖ Several job scripts files available to use for practice
  
- ❖ Edit `job.**.bash.slurm` files as needed, then run the job:
  - you can use gedit
  
- ❖ submit the job(s) i.e.
  - `sbatch job.**.bash.slurm`
  - where `**` represents the job pattern you are going to submit

`job.gethostname.bash.slurm` `job.gethostname.mpi.bash.slurm` `job.gpu.bash.slurm` `job.matlab.bash.slurm`  
`job.namd.bash.slurm`

# Homework

- ∞ Examine the job files and fix potential glitches
- ∞ Submit the jobs.
- ∞ Record the runtime or execution time and send me these timings as well as the job log files

# Overview

- ❖ Accessing the Cluster
- ❖ Maxwell Specifications
- ❖ User Environment
- ❖ Compiling Programs
- ❖ Portable Batch System (PBS)
- ❖ Batch Jobs
- ❖ Interactive Jobs
- ❖ Killing & Suspending Jobs

# First Access Your Account

- ☞ Log into your accounts

- Username or login = hpc\_userX
- Where x = sign in serial number 1 – 47
- Password = **cacds2014**
- Use your web browser
  - Firefox, Chromium or Google chrome

- ☞ Slides could be downloaded from URL below

</share/apps/tutorials/intro2maxwell.pdf>

# Disclaimer

- ∞ Maxwell cluster is for research work under a UH PI
  - Faculty sponsored RESEARCH activity only
- ∞ You will be using a surrogate cluster for this training in place of the Maxwell Cluster
  - You would not be using Maxwell for this training
- ∞ We assuming you are already familiar with the Linux Operating Environment
  - If not try to attend any of the upcoming Linux hands on tutorials offered by CACDS
- ∞ After this training, if you do not already have an account on Maxwell feel free to request one, after authorization from your PI(UH Faculty)
  - <http://www.rcc.uh.edu/cgi-bin/account.cgi> (actual application page)

# Who Can Use Maxwell Cluster?

- ∞ Open to all UH faculty for research purposes
  - Students & postdocs – under UH faculty guidance
  - Collaborators – under UH faculty guidance

# What Happens Next?

## ❖ How to connect to the Maxwell Cluster:

- For UNIX/Linux/Mac users
  - Open up terminal console and use Openssh secure shell client to connect
  - syntax:
    - **ssh username@cusco.hpcc.uh.edu**

Where **username** is your assigned login ID..

- you might be prompted to accept finger print...
  - Just type “yes” followed by Enter key
- For Windows users
  - Install Putty or any other compatible SSH client to connect
  - Putty URL = <http://www.putty.org/>

# Connecting via OpenSSH Client on UNIX based Systems

```
bash-4.2$ ssh -Y jebaluno@cusco.hpcc.uh.edu
The authenticity of host 'cusco.hpcc.uh.edu (129.7.54.9)' can't be established.
RSA key fingerprint is ac:3b:99:d0:6e:84:c0:d8:c9:16:60:d9:fe:e6:c7:71.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'cusco.hpcc.uh.edu,129.7.54.9' (RSA) to the list of known hosts.
```

Use of University of Houston computing and network facilities requires prior authorization. Unauthorized access is prohibited. Usage may be subject to security testing and monitoring. Abuse is subject to criminal prosecution. A complete manual of security policies and procedures is available at <http://www.uh.edu/> in the Administration directory.

jebaluno@cusco.hpcc.uh.edu's password:



MacRumors: G

mac guides and tutori

# Connecting via OpenSSH Client on UNIX based Systems

jebaluno@cusco.hpcc.uh.edu's password:

Last login: Mon Oct 7 11:26:13 2013 from 129-7-249-234.dhcp.uh.edu

On My Computer

Sent Items

UH (University of Houston)

Sent Items

USC (University of South Carolina)

On My Computer

Passwords expire every 90 days.

Deleted Items

On My Computer

Please do not run your code on the login node. It can cause login, compiling, and other problems.

On My Computer

Use module avail to list available software

Use module add <name> to set up environment to use <name>

Example module add matlab

Conflicts

Operated by IT HPC Group | <http://www.rcc.uh.edu>

Help?? Email HPC | [help@hpcc.uh.edu](mailto:help@hpcc.uh.edu)

Quick Reference Doc | </usr/local/doc/maxwellquickref.pdf>

USC (UNIVERSITY OF SOUTH CAROLINA)

Junk E-mail

UH (University of Houston)

USC (University of South Carolina)

[jebaluno@cusco ~]\$ hostname

cusco.hpcc.uh.edu

[jebaluno@cusco ~]\$

AbhishekR@FROST.COM

AbhishekR@FROST.COM

Chapman, Barbara

Ebalunode, Jerry O

ECSS Symposium Oct 15, 10:

XSEDE User News

We've got six GTC Express w

prithvi bs, Chapman, Barbara

Fwd: Reminder: need to get I

Jerry Ebalunode

Fwd: Training update

Jerry Ebalunode

Fwd: Re: Invitation for a talk

Jerry Ebalunode

Confirmation: "OpenACC 2.0"

Nadeem Mohammad

Account details for jebaluno

webmaster@evesopen.com

Wedding Registration for Jerry

Jerry Ebalunode

Upcoming Webinar – Recent

OpenEye Scientific Software

[Kraken-users] NICS stakeho

Gregory D. Peterson

NICS: Kraken emergency mai

XSEDE User News

YESTERDAY

Slides Outline

navigati

Mac

Rec

Ran

Help

search

Go

toolbox

Whe

Rele

Uplc

Spec

Print

Perr

recent p

How

Sno

driv

(0)

Cre

Prim

Mac

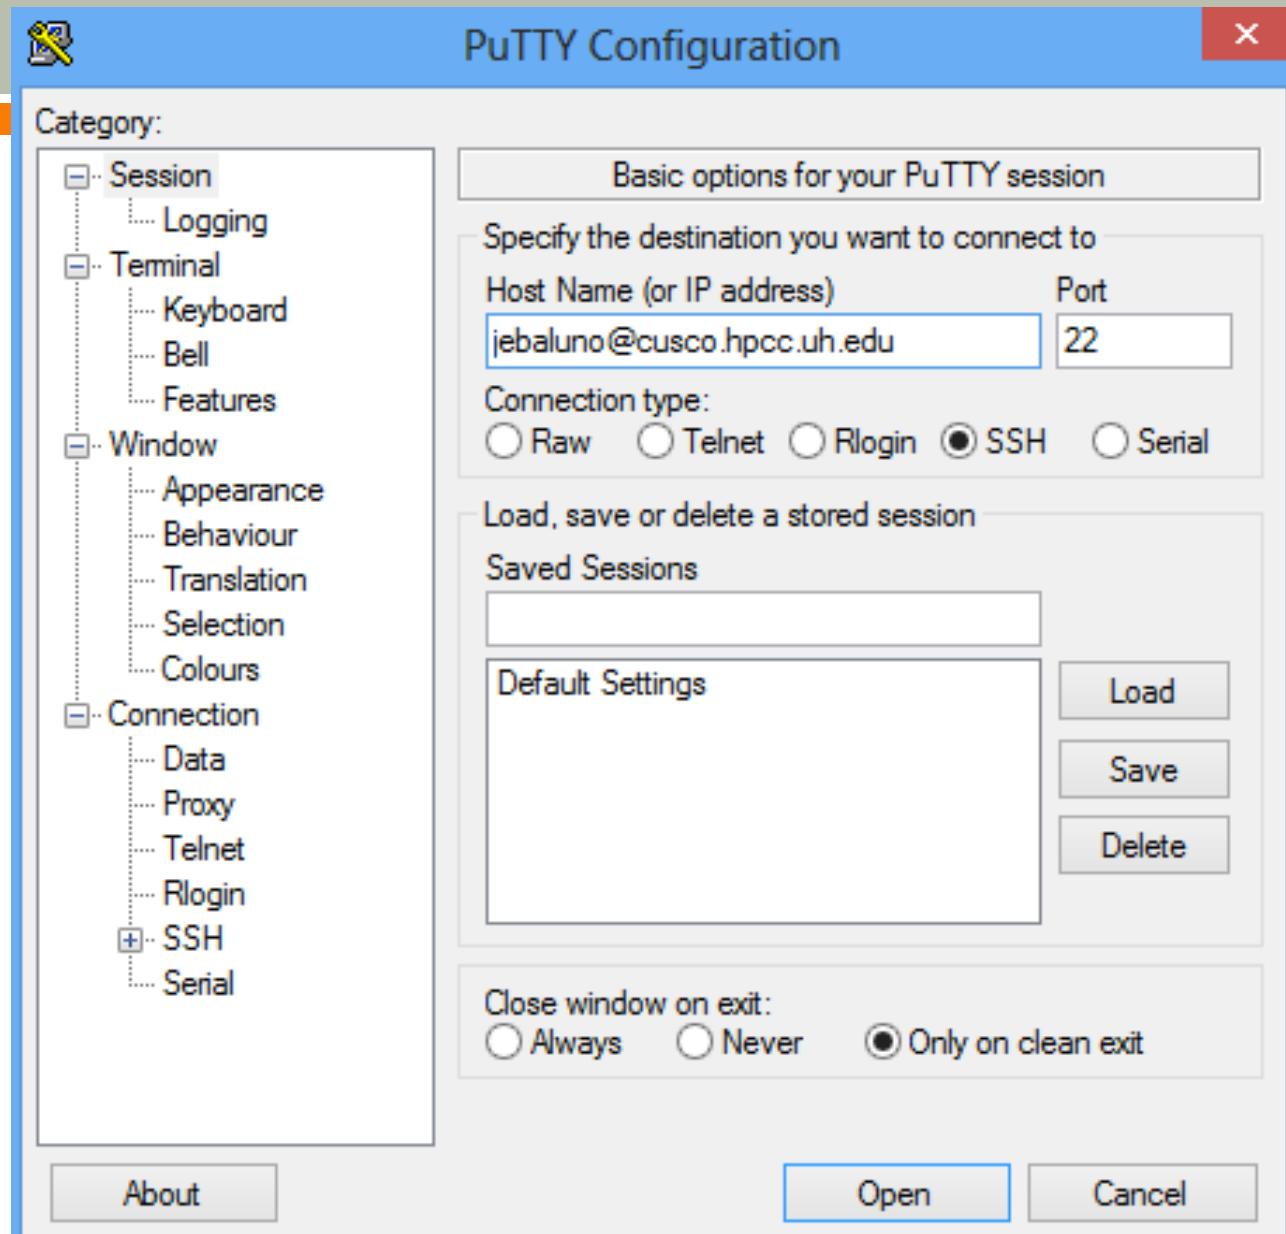
201

exte

prim

chav

# Connecting via Putty on Windows Systems



# Connecting via Putty on Windows Systems

```
Using username "jebaluno".
```

```
Use of University of Houston computing and network facilities requires  
prior authorization. Unauthorized access is prohibited. Usage may be  
subject to security testing and monitoring. Abuse is subject to  
criminal prosecution. A complete manual of security policies and  
procedures is available at http://www.uh.edu/ in the Administration  
directory.
```

```
jebaluno@cusco.hpcc.uh.edu's password: █
```

Note: you will not be able see your password

# Accessing Tutorial Materials

First login into the cluster:

**TYPE AND EXECUTE COMMANDS IN RED!!!**

ssh -Y 10.1.1.1

cd

cp /share/apps/tutorials/intro2maxwell.zip ~

unzip intro2maxwell.zip

cd intro2maxwell

# The Maxwell Cluster

- ∞ Both Shared and Distributed systems
  - single core, multicore (multithreaded) and MPI jobs
- ∞ 434+ compute nodes (3756+ cpu cores)
  - <http://baragon.hpcc.uh.edu/ganglia>
  - 22+ Geforce GTX 570 Nvidia GPUs
  - 8+ Nvidia Tesla C2075 GPUs
- ∞ Operating system: Linux
  - Distribution: Redhat 6.5
- ∞ Gigabit network interconnect
- ∞ Data Storage:
  - Default: 10+ GB Home directory space per user
  - PI's could request for project directory space to host files needed for simulations
    - Justification might be required

# Specifications

Name	Count	Processor Type	Speed GHz	Total CPU Cores	Memory/ core	GPU
cusco (login node)	1	Quad-Core AMD Opteron(tm) Processor 2354	2.20	8	32+ GB	0.00
compute-X-Y	183	Quad-Core AMD Opteron(tm) Processor 2354	2.20	8	16+ GB	0.00
compute-X-Y	25	Quad-Core AMD Opteron(tm) Processor 8356	2.30	32	64+ GB	0.00
compute-X-Y	39	Quad-Core AMD Opteron(tm) Processor 2376	2.30	8	16.3+ GB	0.00
compute-X-Y	112	Dual Core AMD Opteron(tm) Processor 275	2.20	4	8+ GB	0.00
compute-X-Y	28	AMD Opteron(TM) Processor 6274	2.50	32	65+ GB	0.00
compute-X-Y	22	Intel(R) Xeon(R) CPU 5160	3.00	4	32+ GB	GeForce GTX 570
compute-X-Y	8	Intel(R) Xeon(R) CPU 5160	3.00	4	32+ GB	Tesla C2075

X = Rack Number

Y = Node Number

# HPC Software Available on Maxwell

## ∞ Proprietary & Open-Source Software:

- Examples:
  - Amber, Autodock, Bwa, NwChem, Espresso, Orca, R, Octave, Matlab, Gromacs, Lammgs, NAMD, Turbomole, VASP, Rosetta, R, NAG, MKL , ASE, GSL, OpenMPI, Intel, Gnu and PGI compilers ....
  - and much more...
- Proprietary software may be limited to licensed users e.g. Fluent, VASP
- Open-source software is accessible to all

## ∞ PI could request for installation of additional software

- PI's licensed propriety application can be installed too

## ∞ Software environment management via “**modules**”

- Allows software to be accessed via an assigned module
- Dynamic modification of a user's environment

# Computing Environment Setup Using Modules

- ❖ The module system is used to make software and related settings available easily
  - software environments can be loaded and unloaded dynamically

- ❖ To get a list of available software:

`module avail`

(also useful: `module list`)

- ❖ To clear-out all added modules:

`module purge`

- ❖ To clear-out all specific modules:

`module rm module1 module2 ....`

- ❖ To use specific software every time you log-in, place the

`module add` command in your `~/.bashrc` file i.e:

`echo "module add intel openmpi" >> ~/.bashrc`

- ❖ Now add intel compiler and MPI runtime (Openmpi version) to your environment)

`module add openmpi`

# Compiling Programs

- ❖ CPU programs (compiled to execute on central processing unit)
  - Serial Programs
  - Parallel Programs
- ❖ GPU programs (compiled to execute on graphics processing unit)
  - GPGPU Programs

# Compiling Programs: CPU Serial Program

## Example 1 Program in file: gethostname.c

```
#include <stdio.h>
#include <sys/utsname.h>
int main ( )
{
    struct utsname uts;
    uname (&uts);
    printf ("Process on node %s.\n",
    uts.nodename);
    return 0;
}
```

- To compile:
  - 1. choose compiler  
`module add intel`
  - 2. Compile  
`icc gethostname.c -o gethostname.exe`

# Compiling Programs: CPU Parallel Program

## Programs using Message Passing Interface (MPI)

Example: 1 Program in file : gethostname.mpi.c

```
#include <stdio.h>
#include <sys/utsname.h>
#include <mpi.h>
int main (int argc, char *argv[]){
    struct utsname uts;
    int rank;
    MPI_Init (&argc, &argv);
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    uname (&uts);
    printf ("Process %d on node %s.\n", rank, uts.nodename);
    MPI_Finalize ();
    return 0;
}
```

- To compile:
  - 1. choose compiler  
`module add openmpi`
  - 2. Compile  
`mpicc gethostname.mpi.c -o gethostname.mpi.exe`

# Compiling Programs: GPGPU Parallel Program

## Example 1 Program in file: sortkeys\_basic\_thrust.cu

```
#include <thrust/host_vector.h>
#include <thrust/device_vector.h>
#include <thrust/sort.h>
int main(void)
{
    // generate 16 Million random numbers on the host
    thrust::host_vector<int> h_vec(1 << 24);
    thrust::generate(h_vec.begin(), h_vec.end(), rand);
    thrust::device_vector<int> d_vec= h_vec; // transfer data to the device
    thrust::sort(d_vec.begin(), d_vec.end()); // sort data on the device (805 Mkeys/sec on GeForce GTX 480)
    thrust::copy(d_vec.begin(), d_vec.end(), h_vec.begin()); // transfer data back to host
    return 0;
}
```

- To compile:
  - 1. choose compiler

```
module add cuda-toolkit
```

```
2. Compile name = nvcc
```

```
nvcc sortkeys_basic_thrust.cu -o sortkeys_basic_thrust.exe
```

# The Queue System (PBS)



- ❖ Portable Batch System (aka PBS) is the resource management service used on the Maxwell Cluster
- ❖ Enables you to make more efficient use of your time through scripting computational tasks
- ❖ PBS takes care of running these tasks and returning the results
- ❖ If the cluster is full, PBS holds your tasks and runs them when the resources are available
- ❖ PBS ensures fair sharing of cluster resources (policy enforcement)
  - PBS ensures optimal/efficient use of available resources

# PBS Commands

## Command(s)

## Description

**qstat**

Check status of all jobs

**qsub ./myjob**

Submit batch jobs

**qstat -u \$USER**

Check status of just your jobs

**qsub -I**

Submit an interactive job

**qhold jobID**

i.e. qhold 12345

Put a job on hold (before it starts)

**qrsl jobID**

Release a job from hold status

**qdel #jobID**

Delete a job, running or not



# Using qsub with a Job Script

- ❖ Syntax/Format:

- **qsub [qsub params] [job script file]**

- ❖ Examples:

- **qsub myjob**

# PBS Batch Job Script File

- ❖ Very simple example:
- ❖ Serial job requesting 1 CPU, 2 GB memory, 60 hours of walltime

```
#!/bin/bash  
#PBS -N jobname  
#PBS -l mem=2gb,walltime=60:00:00  
#PBS -j y
```

```
cd $PBS_O_WORKDIR
```

```
./my.program.exe [arguments for my program]
```

# Using Parallel Environment to Manage Memory and CPU Core Availability

- ☞ `#PBS -l nodes=16:ppn=1` --use only one CPU core per node and use a total of 16 processes.
- ☞ `#PBS -l nodes=8:ppn=2` --use only 2 CPU cores per node and use a total of 16 processes
- ☞ `#PBS -l nodes=4:ppn=4` --use only 4 CPU cores per node and use a total of 16 processes
- ☞ `#PBS -l nodes=8:ppn=8` --use only 8 CPU cores per node and use a total of 64 processes

# qsub Notification Parameters

## ∞ #PBS -M [e-mail address]

- e-mail address can be a list of email addresses
- separated by commas.

## ∞ #PBS -m bea or -m be or -m e

- send an email when the job **b**egins, **e**nds, or is **a**borted.

**Example:**

```
#!/bin/bash
#PBS -N jobname
#PBS -j y
#PBS -M jerry@uh.edu
#PBS -m abe
```

```
cd $PBS_O_WORKDIR
```

```
./my.program.exe [arguments for my program]
```

# Sample Notification Email

**PBS JOB 3687125.cusco.hpcc.uh.edu**

root

**Sent:** Monday, October 7, 2013 at 6:27 AM

**To:** jebalunode@uh.edu

PBS Job Id: 3687125.cusco.hpcc.uh.edu

Job Name: waterint-ACC-H17

Exec host: compute-7-25/31+compute-7-25/30+compute-7-25/29+compute-7-25/28+compute-7-25/27+compute-7-25/26+compute-7-25/25+compute-7-25/24+compute-7-25/23+compute-7-25/22+compute-7-25/21+compute-7-25/20+compute-7-25/19+compute-7-25/18+compute-7-25/17+compute-7-25/16+compute-7-25/15+compute-7-25/14+compute-7-25/13+compute-7-25/12+compute-7-25/11+compute-7-25/10+compute-7-25/9+compute-7-25/8+compute-7-25/7+compute-7-25/6+compute-7-25/5+compute-7-25/4+compute-7-25/3+compute-7-25/2+compute-7-25/1+compute-7-25/0

Execution terminated

Exit\_status=0

resources\_used.cput=08:46:15

resources\_used.mem=352472kb

resources\_used.vmem=35667756kb

resources\_used.walltime=00:33:07

# Handling Output Files

- ❖ When a job is started it takes its job name from the script file that was submitted.
  - The standard output and error output are sent to files named jobname.o987349 and jobname.e987349 (in your job directory )
- ❖ The following parameters modify this behavior:
  - **-e** [path] standard error output file.
  - **-o** [path] standard output file
  - **-j y** merge the error and standard output
- ❖ **#PBS -N** jobname name to be used for the job.

# Using qsub with a Job Script

## ❖ Syntax/Format:

- `qsub [qsub params] [script file]`

## ❖ More examples on using parameters at the command line:

- `qsub -N test1 myjob.sh`
- `qsub -l nodes=2:ppn=2 -o test2.out -N test2 myjob.sh`
- `qsub -l walltime=02:00:00,mem=4gb myjob.sh`
- `qsub -l nodes=2:ppn=32,mem=64gb myjob.sh`
- `qsub -q gpu -l nodes=1:ppn=4 my_gpu_job.sh`

## ❖ To see all available qsub options, just run `qsub --help`

# Using an Interactive Job

- ∞ Special kind of batch job
- ∞ Useful for debugging applications, short tests, or for computational steering

The diagram illustrates the usage of interactive job options. It features two orange callout boxes at the top: 'Capital “i”' on the left and 'Small L “l”' on the right. Two arrows point downwards from these labels to corresponding command examples. One orange arrow points from 'Capital “i”' to the first command, and one brown arrow points from 'Small L “l”' to the second command. The remaining three commands are aligned horizontally below the second one.

```
qsub -I
qsub -I -l walltime=2:00:00
qsub -I -l walltime=2:00:00,nodes=2:ppn=8
qsub -I -l walltime=2:00:00 -q gpu
```

# Job Queues

Queue Name	Description
short	Queue to run short jobs ( =< 4 hours)
medium	Queue to run medium jobs ( 1 week allowed)
long	Queue to run very long jobs (2 weeks allowed)
gpu	Queue to run gpu jobs on GTX 570 equipped nodes (2 weeks allowed)
gpu-tesla	Queue to run gpu jobs on Testla C2075 equipped nodes (2 weeks allowed)

USE “qstat -q” to probe the queues installed

Note gpu-tesla and gpu queue are not available/functional on the surrogate cluster

# PBS Script for Serial Program

## 1 CPU

### Example of a Serial Job Script

filename: job.gethostname.bash

```
#!/bin/bash
#PBS -N gethost
#PBS -l mem=8gb,walltime=00:25:00
#PBS -l nodes=1:ppn=1
#PBS -M JohnDoe@gmail.com
#PBS -m bea

cd $PBS_O_WORKDIR

### Run gethostname

time ./gethostname.exe
```

To submit the job run:

**qsub job.gethostname.bash**

# PBS Script for MPI Program

```
#!/bin/bash
#PBS -N gethostmpi
#PBS -l nodes=2:ppn=8,pmem=1gb
#PBS -S /bin/bash
#PBS -l walltime=00:05:00
#PBS -M monkeybrain@sc.edu
#PBS -m bea

cd $PBS_O_WORKDIR
##set up your environment
source /etc/profile.d/modules.sh

module add openmpi

mpirun -np 16 ./gethostname.mpi.exe > output.from.txt
```

See job filename: job.gethostname.mpi.bash

To submit the job run:

qsub job.gethostname.mpi.bash

# PBS Script for GPGPU Program

## Example GPU Job Script

```
#!/bin/bash
#PBS -q gpu
#PBS -N sortrandomkeys
#PBS -l walltime=00:05:00,nodes=1:ppn=4
#PBS -M monkeybrain@gmail.com
#PBS -m bea

cd $PBS_O_WORKDIR
##set up your environment
source /etc/profile.d/modules.sh
module add cuda-toolkit
time ./sortkeys_basic_thrust.exe
```

Give me a node with  
this kind of gpu

Load CUDA  
environment

See job filename: job.gpu.bash or job.gpu.sort.bash

# PBS Script for Executing Matlab script Using Proprietary Matlab runtime (sample I)

```
#!/bin/bash
#PBS -N matlabjob
#PBS -l nodes=1:ppn=1,pmem=1gb
#PBS -S /bin/bash
#PBS -l walltime=00:05:00

cd $PBS_O_WORKDIR

##set up your environment
source /etc/profile.d/modules.sh
module add matlab

matlab << EOF
a = 10; b = 20; c = 30;
d = sqrt((a + b + c)/pi)
exit
EOF
```

See job filename: job.matlab1.bash

# PBS Script for Executing Matlab script Using Proprietary Matlab runtime (sample II)

```
#!/bin/bash
#PBS -N matlabjob
#PBS -S /bin/bash
#PBS -I walltime=20:05:00,pmem=1gb,nodes=1:ppn=4
#PBS -M monkeybrain@uh.edu
#PBS -m bea
##set up your environment
cd $PBS_O_WORKDIR
source /etc/profile.d/modules.sh
module add matlab

## run matlab program compute pseudo inverse for 100 matrices of size 400x400
time matlab < matrix_inversion100_matlab.m
```

See job filename: job.matlab2.bash

# PBS Script for Executing R script

```
#!/bin/bash
#PBS -N R-job
#PBS -S /bin/bash
#PBS -I walltime=0:05:00,pmem=1gb,nodes=1:ppn=1
#PBS -M monkeybrain@uh.edu
#PBS -m bea
##set up your environment
cd $PBS_O_WORKDIR
source /etc/profile.d/modules.sh
module add R

## run R program
R -vanilla < sample.r
```

See job filename: job.R.bash

# PBS Script for NAMD2 Job Parallel MPI

```
#!/bin/bash
#PBS -l nodes=1:ppn=4
#PBS -l walltime=00:05:00
#PBS -l pmem=1gb
#PBS -N namd
#PBS -V

cd $PBS_O_WORKDIR

module load namd

mpirun -v namd2 apoa1.namd
```

See job filename: job.namd mpi.bash

# Killing Jobs

❖ One, or a few jobs:

- `qdel [jobID] [jobID] [jobID] ...`

❖ Kill all of your jobs:

- `qselect -u $USER | xargs qdel`

❖ Kill all of your queued jobs:

- `qselect -u $USER -s Q | xargs qdel`

❖ Kill all of your running jobs:

- `qselect -u $USER -s R | xargs qdel`

# Try An Example Job

- ❖ Several job scripts files available to use for practice
- ❖ Edit job.\*\*.bash files as needed, then run the job:
  - you can use gedit
- ❖ submit the job(s) i.e.
  - qsub job.\*\*.bash
  - where \*\* represents the job pattern you are going to submit

job.gethostname.bash job.gethostname.mpi.bash job.gpu.bash job.matlab.bash job.namd.bash

# Homework

- ∞ Examine the job files and fix potential glitches
- ∞ Submit the jobs.
- ∞ Record the runtime or execution time and send me these timings as well as the job log files

# Support

- ∞ CACDS Support
  - ∞ Visit → <http://support.cacds.uh.edu>
    - To Open A Support Ticket
- ∞ Jerry Ebalunode
  - jebalunode@uh.edu
  - PGH 231
  - Office hours: 10-12pm Wednesday
  - 713-743-1755